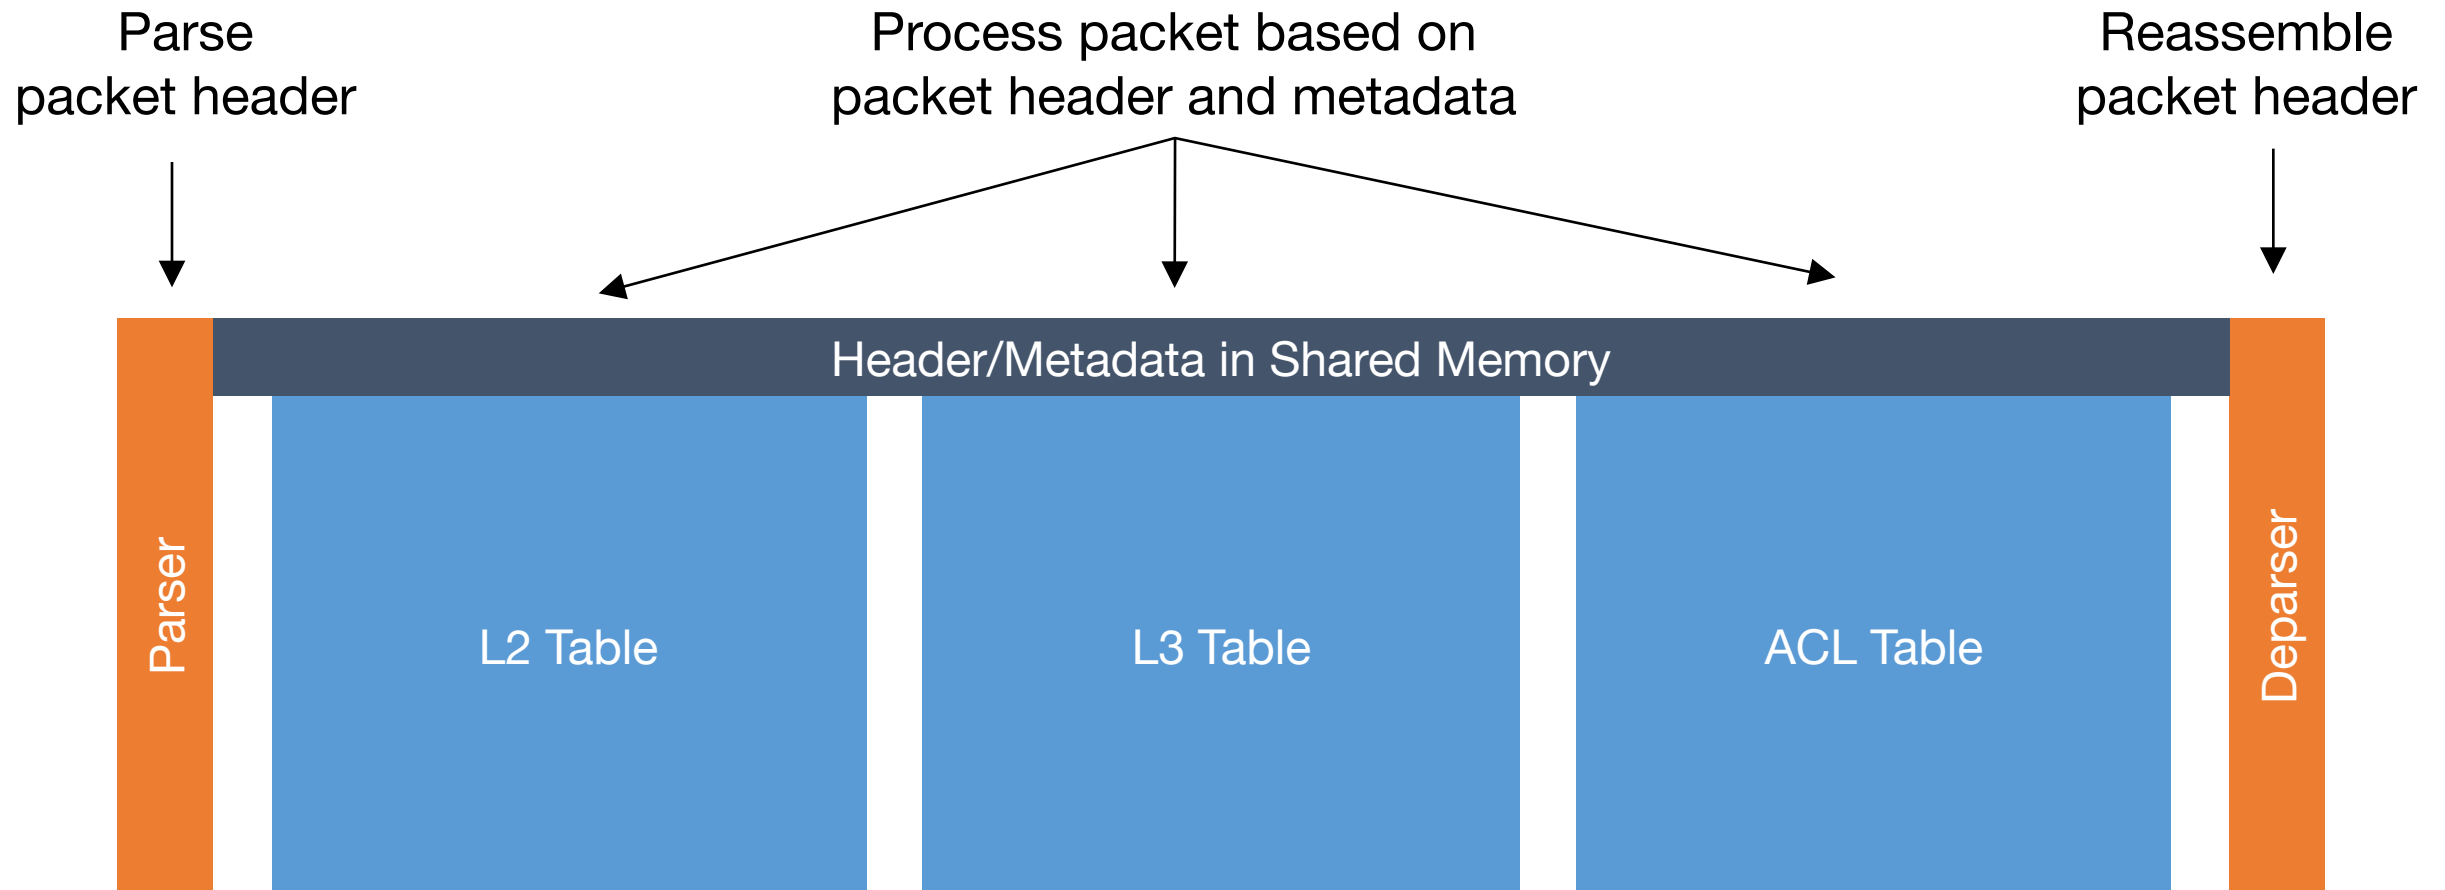


P4 Tutorial

Xin Jin

xinjin@cs.jhu.edu


Traditional Switch: Fixed Packet Processing



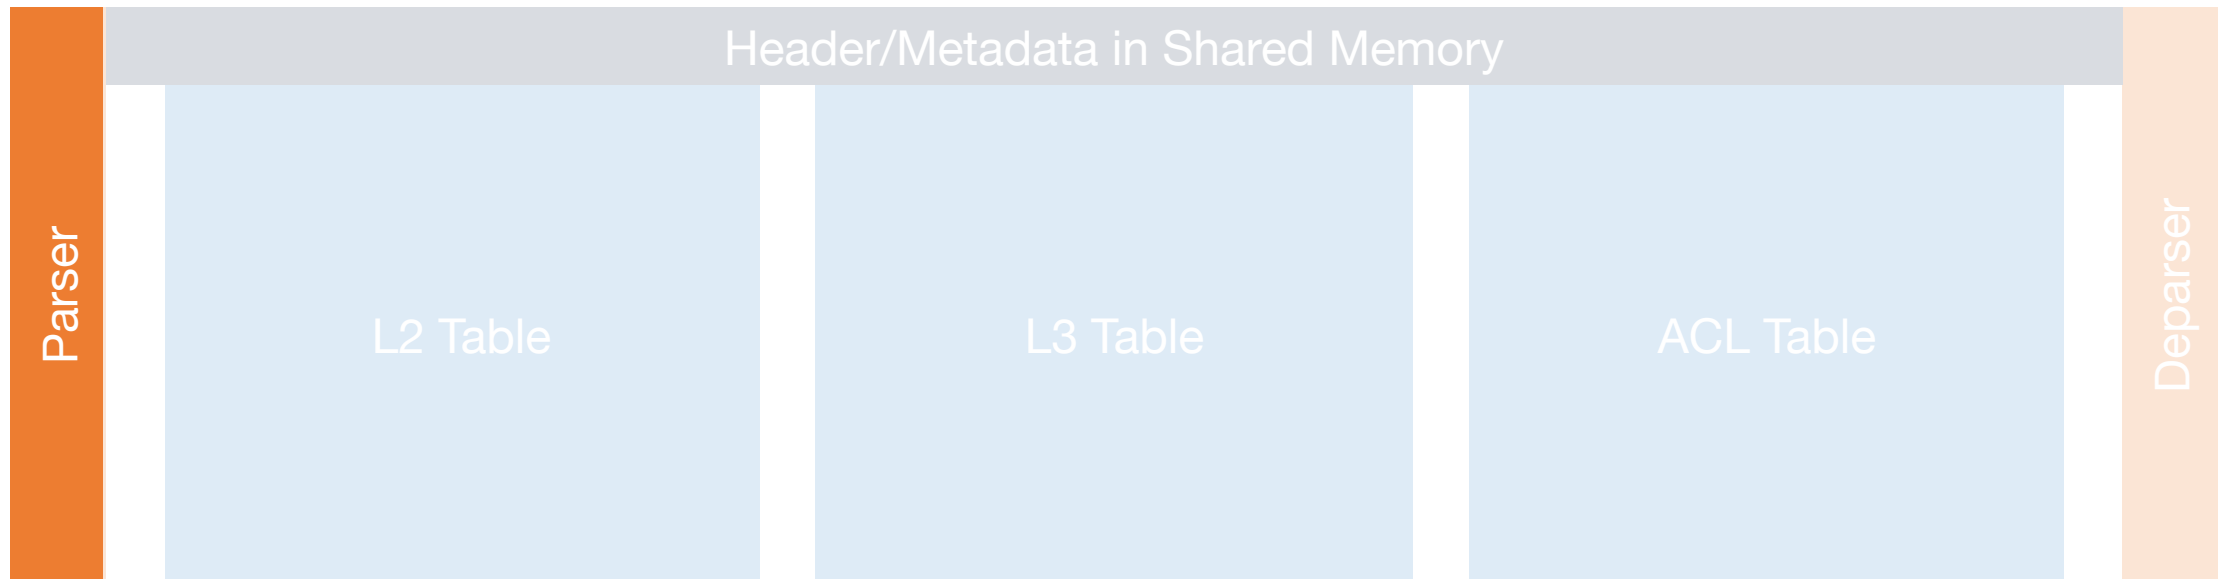

Traditional Switch:

Fixed Packet Processing

Packet format



Parser graph



Traditional Switch:

Fixed Packet Processing

Packet format


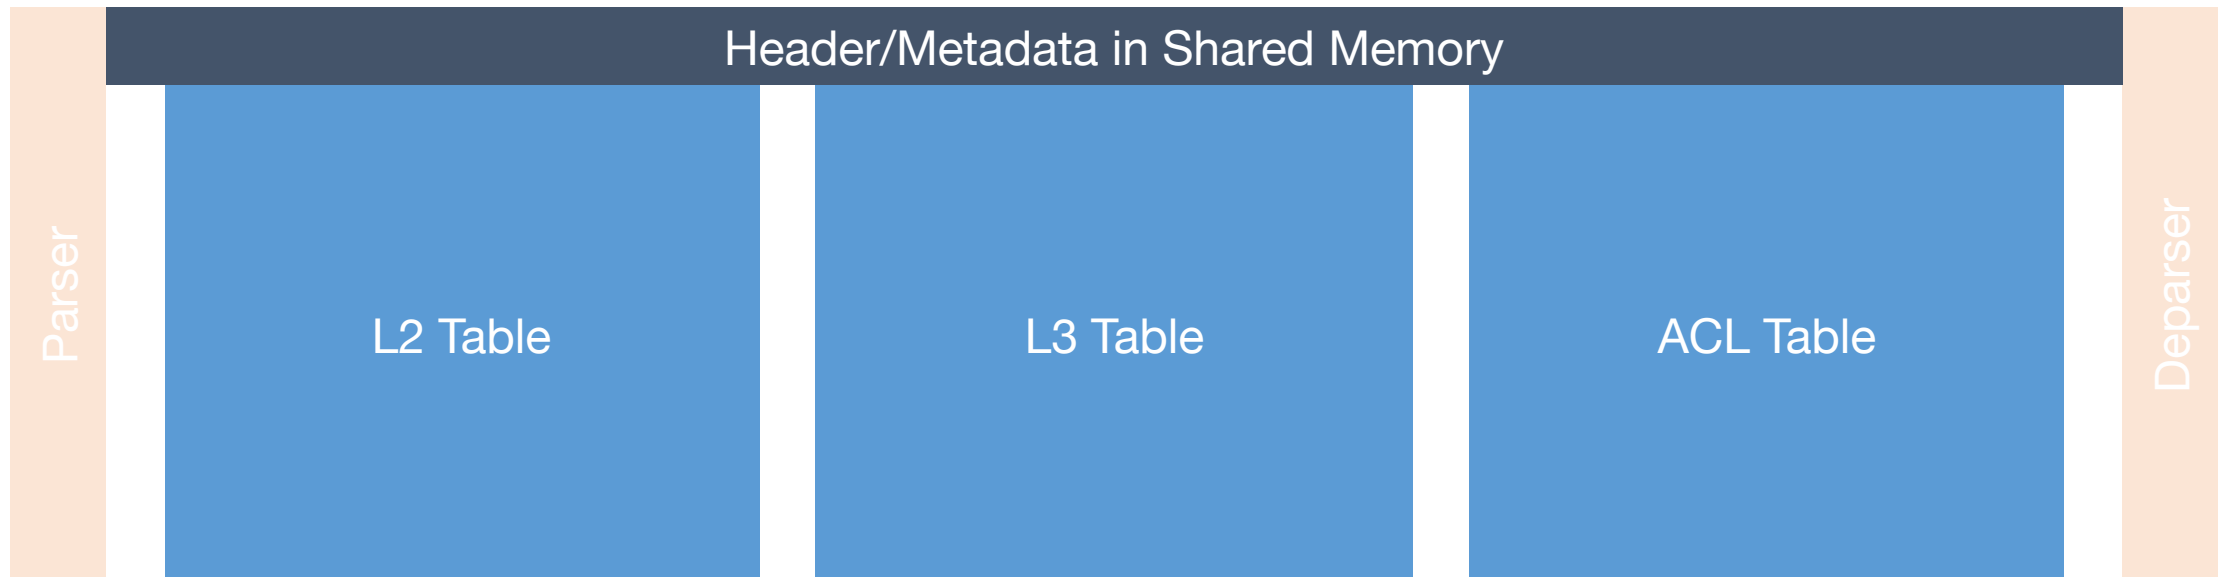


Table graph



Traditional Switch:

Fixed Packet Processing

Packet format


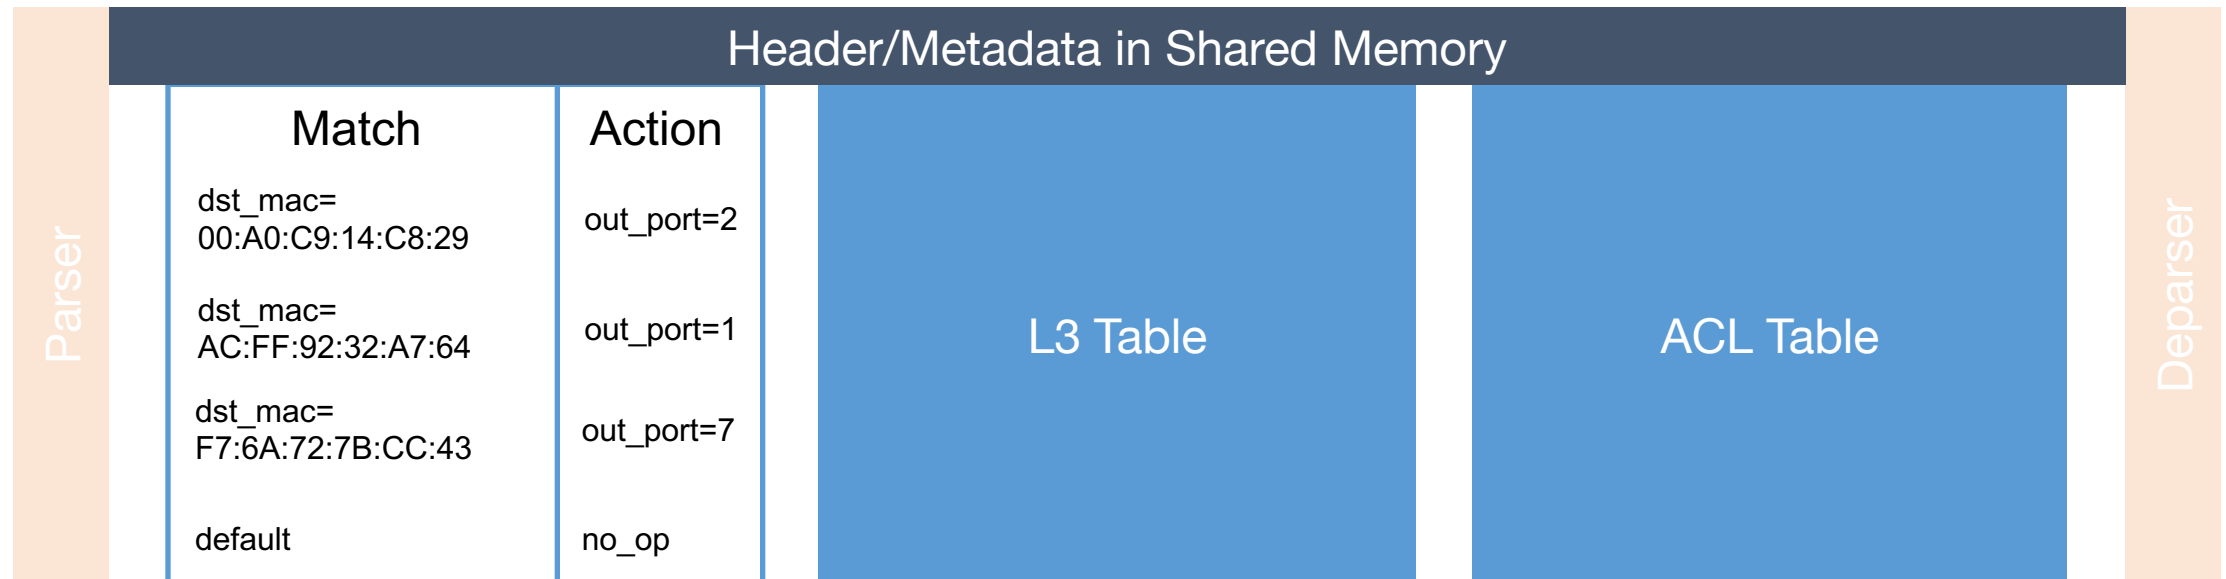


Table graph



Traditional Switch:

Fixed Packet Processing

Packet format


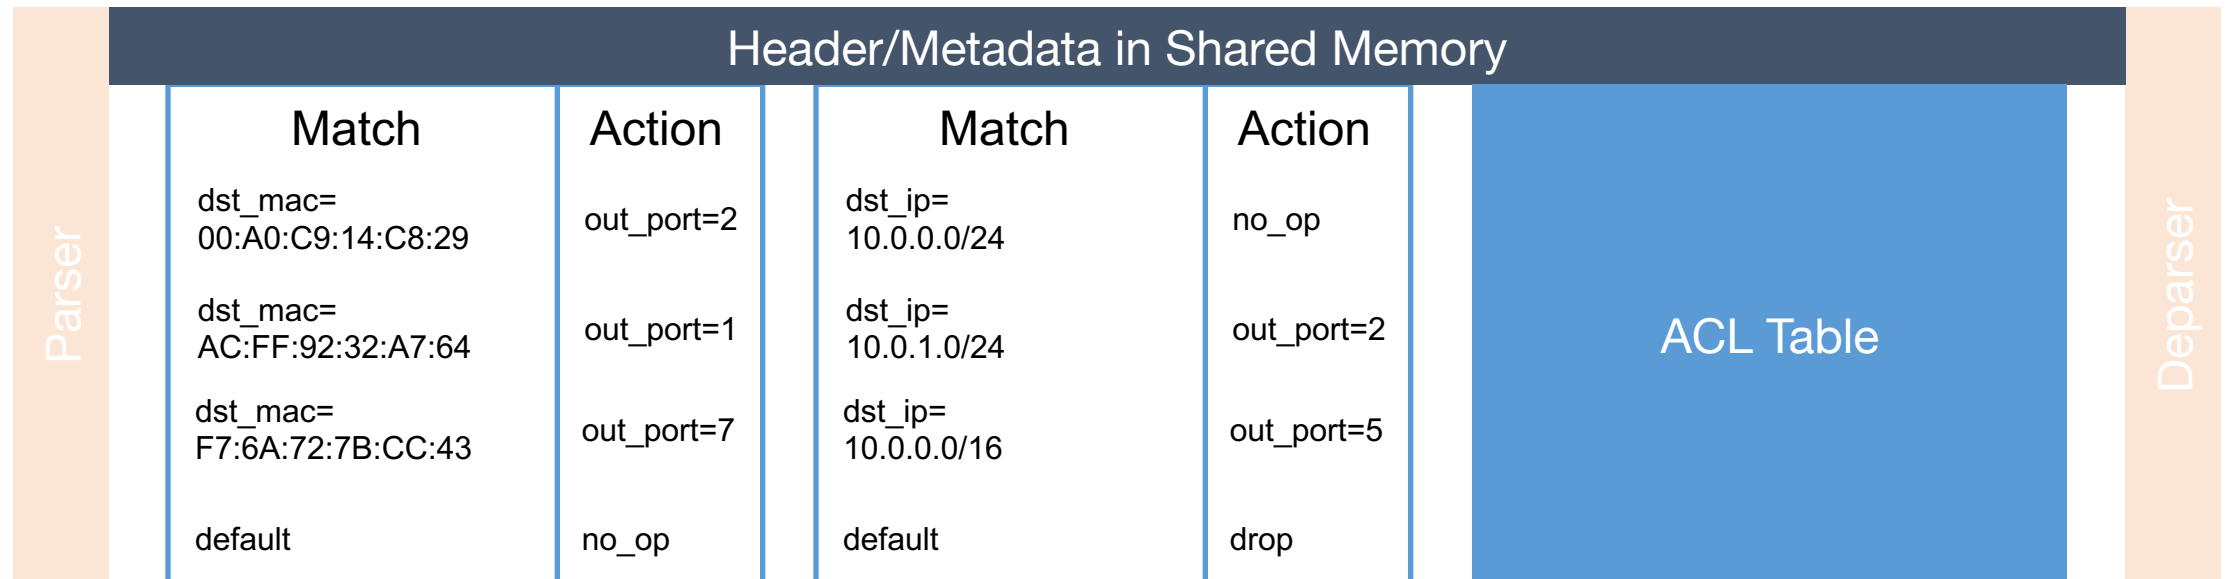


Table graph



Traditional Switch: Fixed Packet Processing

Packet format




Table graph



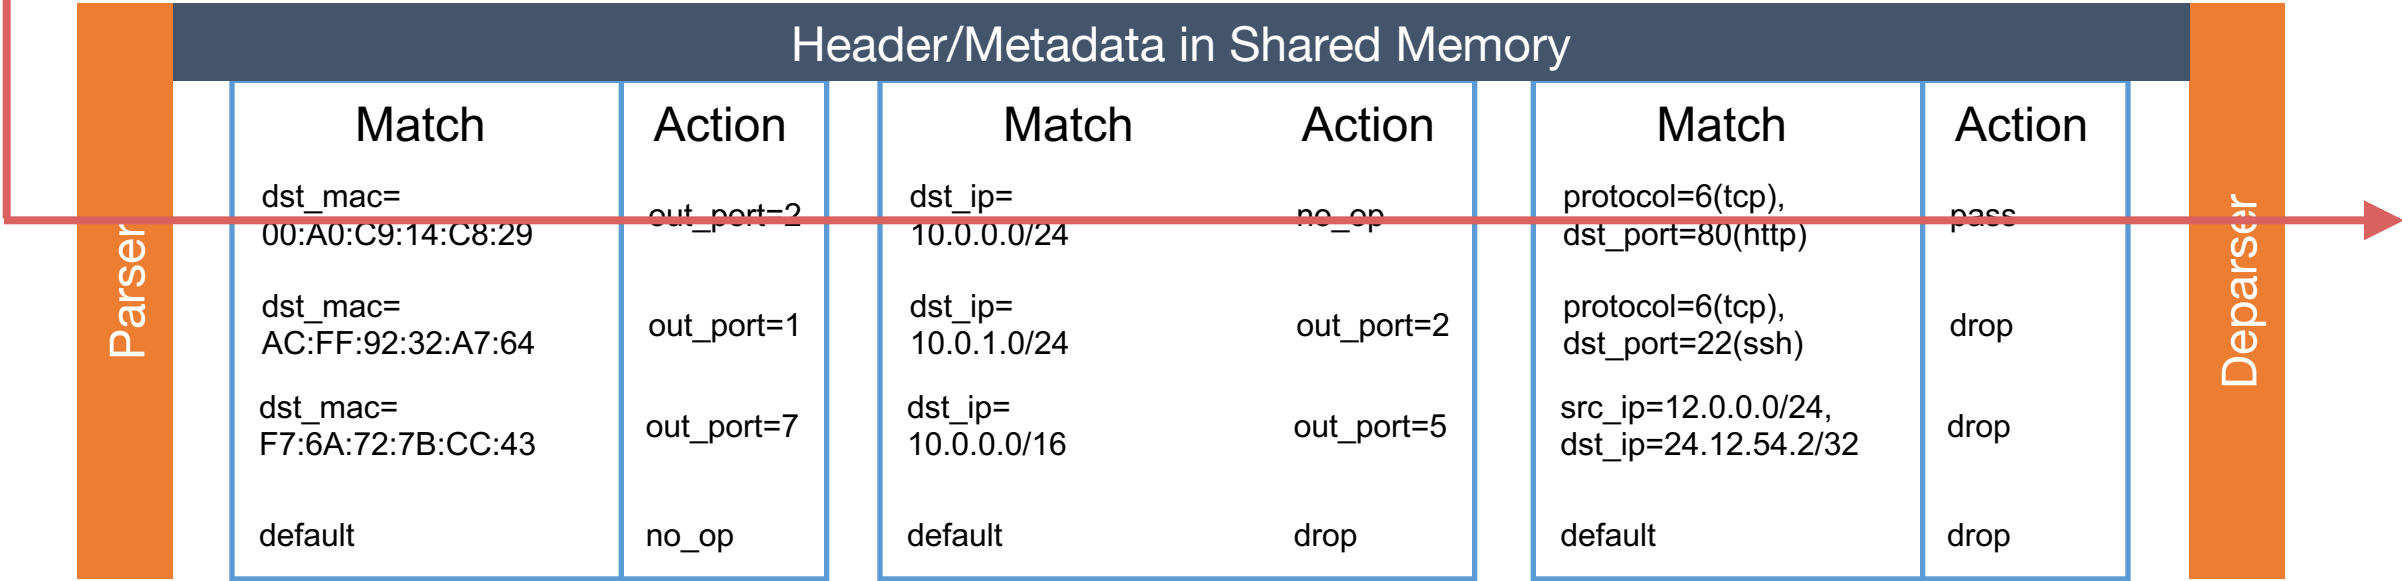
Header/Metadata in Shared Memory							
Parser	Match	Action	Match	Action	Match	Action	Deparser
	dst_mac= 00:A0:C9:14:C8:29	out_port=2	dst_ip= 10.0.0.0/24	no_op	protocol=6(tcp), dst_port=80(http)	pass	
	dst_mac= AC:FF:92:32:A7:64	out_port=1	dst_ip= 10.0.1.0/24	out_port=2	protocol=6(tcp), dst_port=22(ssh)	drop	
	dst_mac= F7:6A:72:7B:CC:43	out_port=7	dst_ip= 10.0.0.0/16	out_port=5	src_ip=12.0.0.0/24, dst_ip=24.12.54.2/32	drop	
	default	no_op	default	drop	default	drop	

Traditional Switch:

Fixed Packet Processing

Example Packet

dst_mac=00:A0:C9:14:C8:29	dst_ip=10.0.0.1, protocol=6	dst_port=80	Payload
---------------------------	--------------------------------	-------------	---------

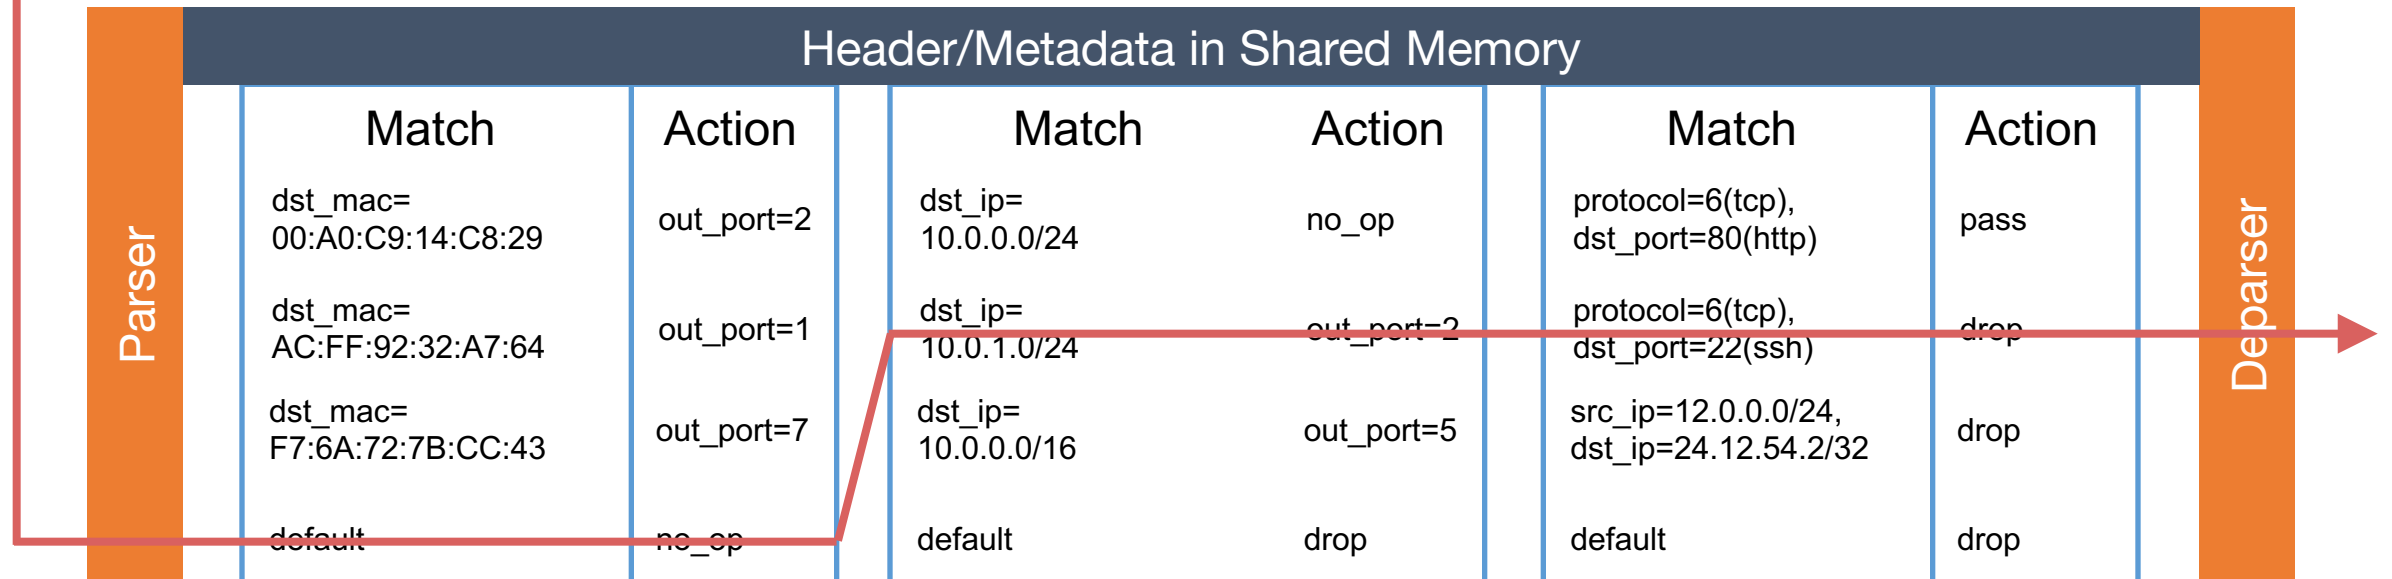


Traditional Switch:

Fixed Packet Processing

Example Packet

dst_mac=11:CC:C9:33:85:2A	dst_ip=10.0.1.1, protocol=6	dst_port=22	Payload
---------------------------	--------------------------------	-------------	---------



Traditional Switch:

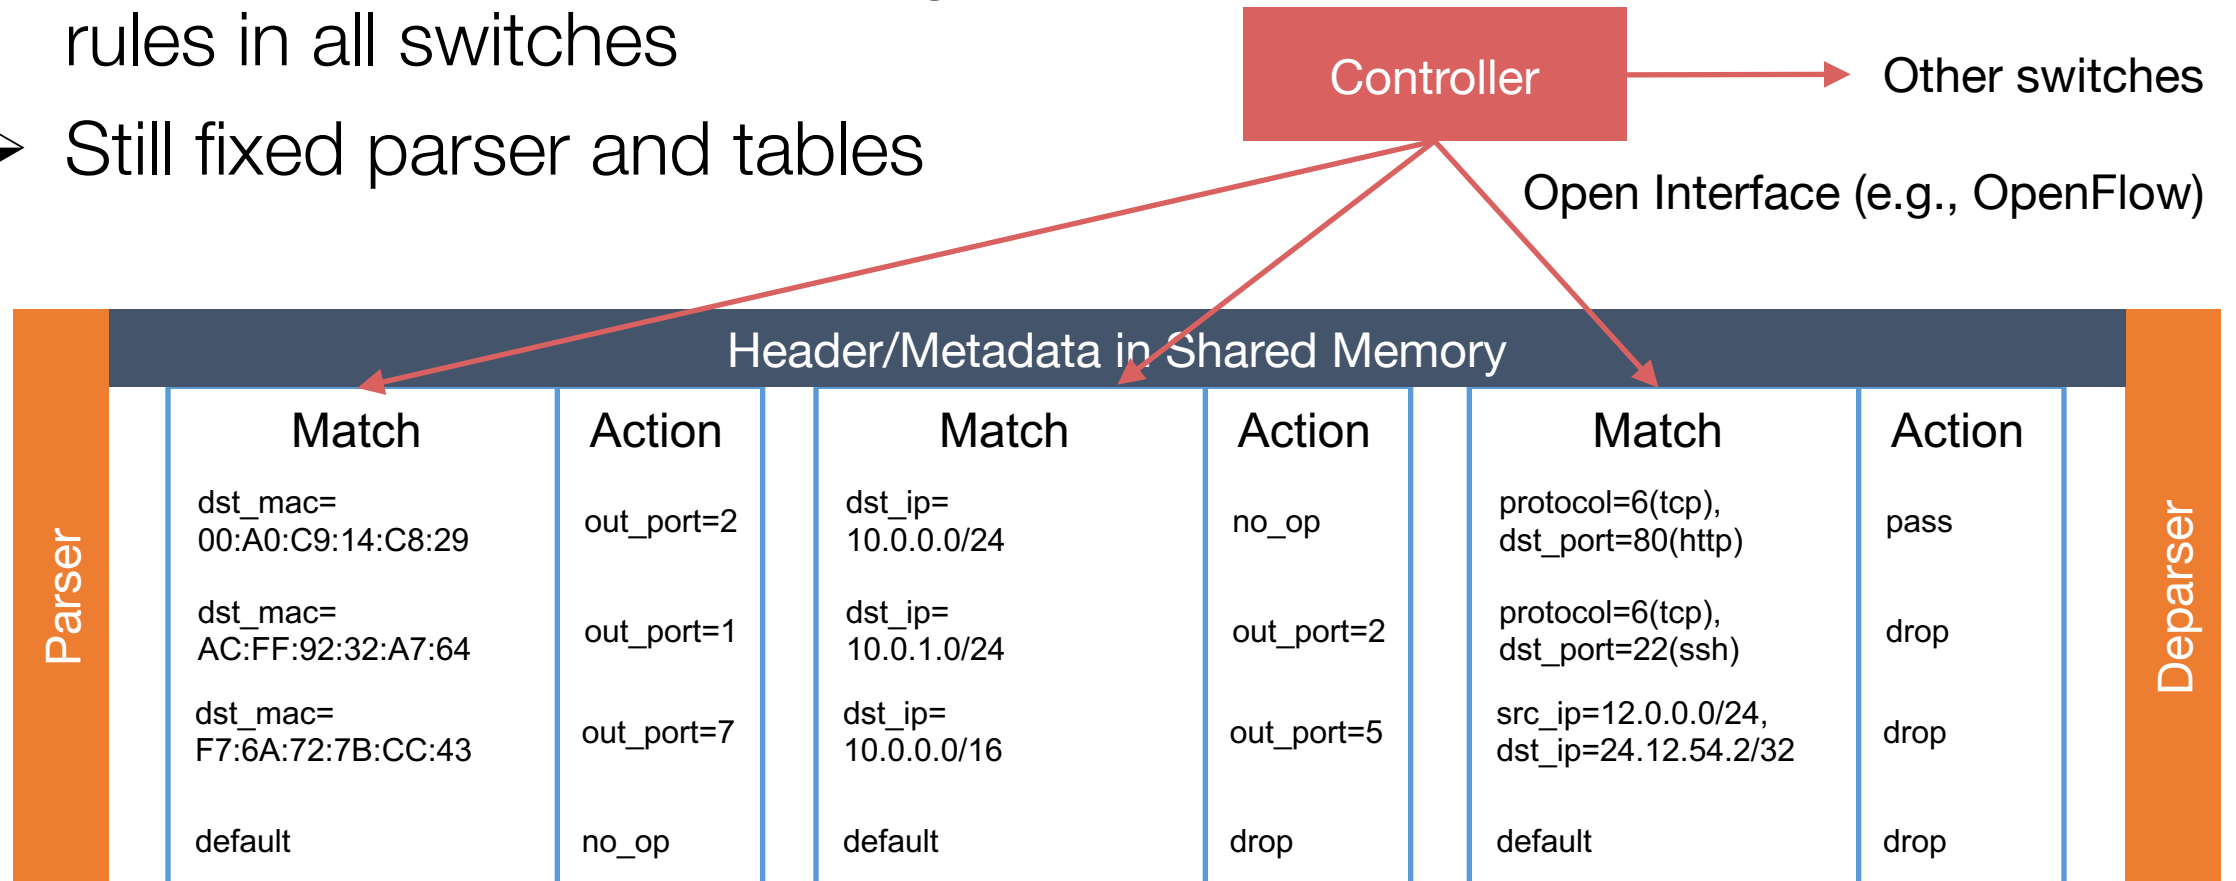
Fixed Packet Processing

- Fixed parser
- Fixed match-action tables
 - Rules in tables are generated by (proprietary) distributed network protocols or configured with (proprietary) switch management interface

Parser	Header/Metadata in Shared Memory								Deparser				
	Match		Action		Match		Action			Match		Action	
	dst_mac=00:A0:C9:14:C8:29		out_port=2		dst_ip=10.0.0.0/24		no_op			protocol=6(tcp), dst_port=80(http)		pass	
	dst_mac=AC:FF:92:32:A7:64		out_port=1		dst_ip=10.0.1.0/24		out_port=2			protocol=6(tcp), dst_port=22(ssh)		drop	
	dst_mac=F7:6A:72:7B:CC:43		out_port=7		dst_ip=10.0.0.0/16		out_port=5			src_ip=12.0.0.0/24, dst_ip=24.12.54.2/32		drop	
	default		no_op		default		drop		default		drop		

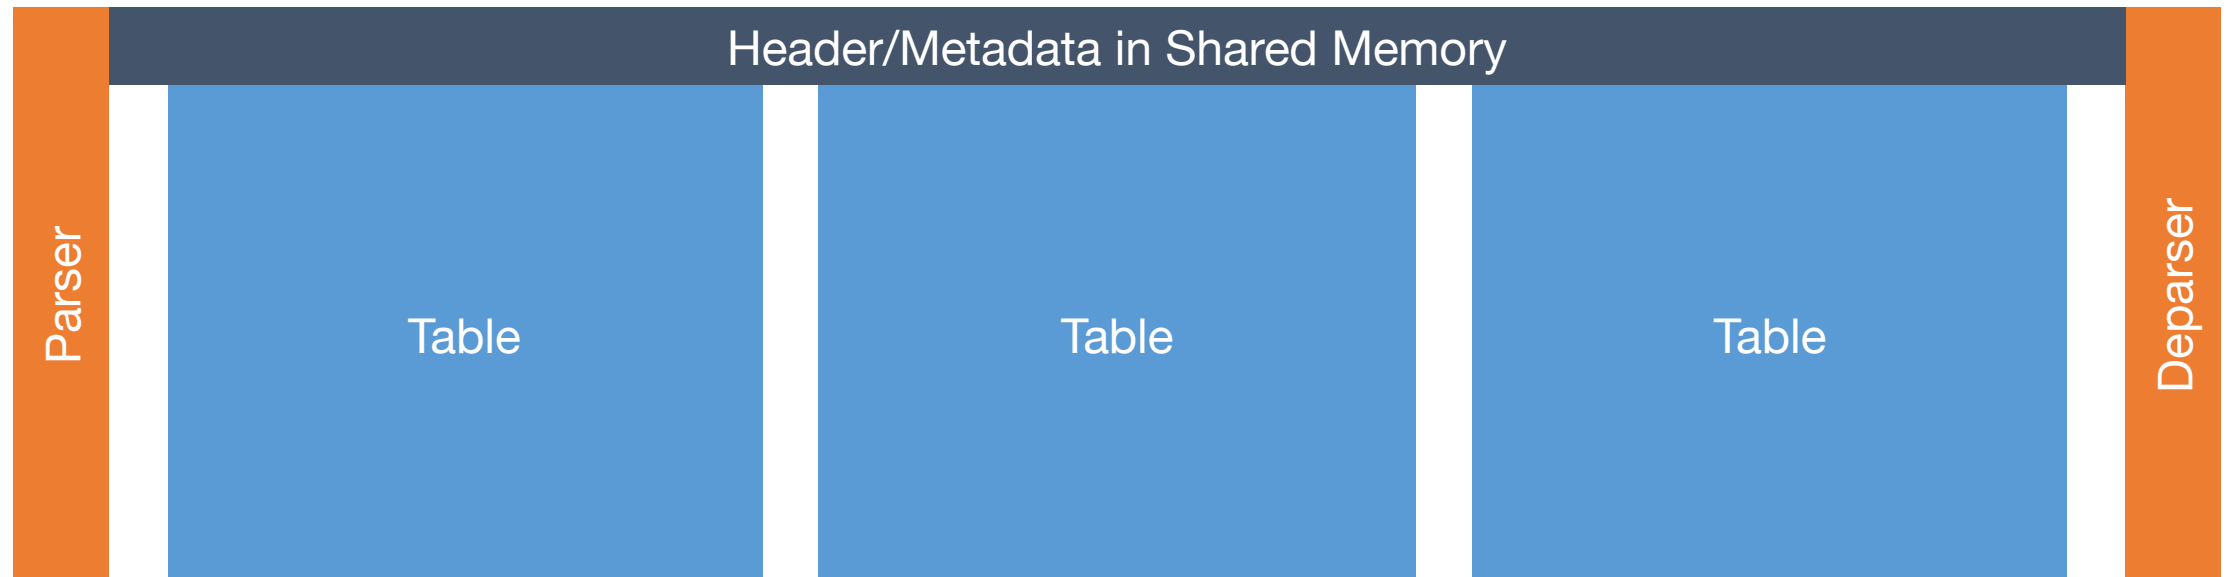
Software-Defined Networking: Centralized Control of Switches

- Centralized controller configures rules in all switches
- Still fixed parser and tables



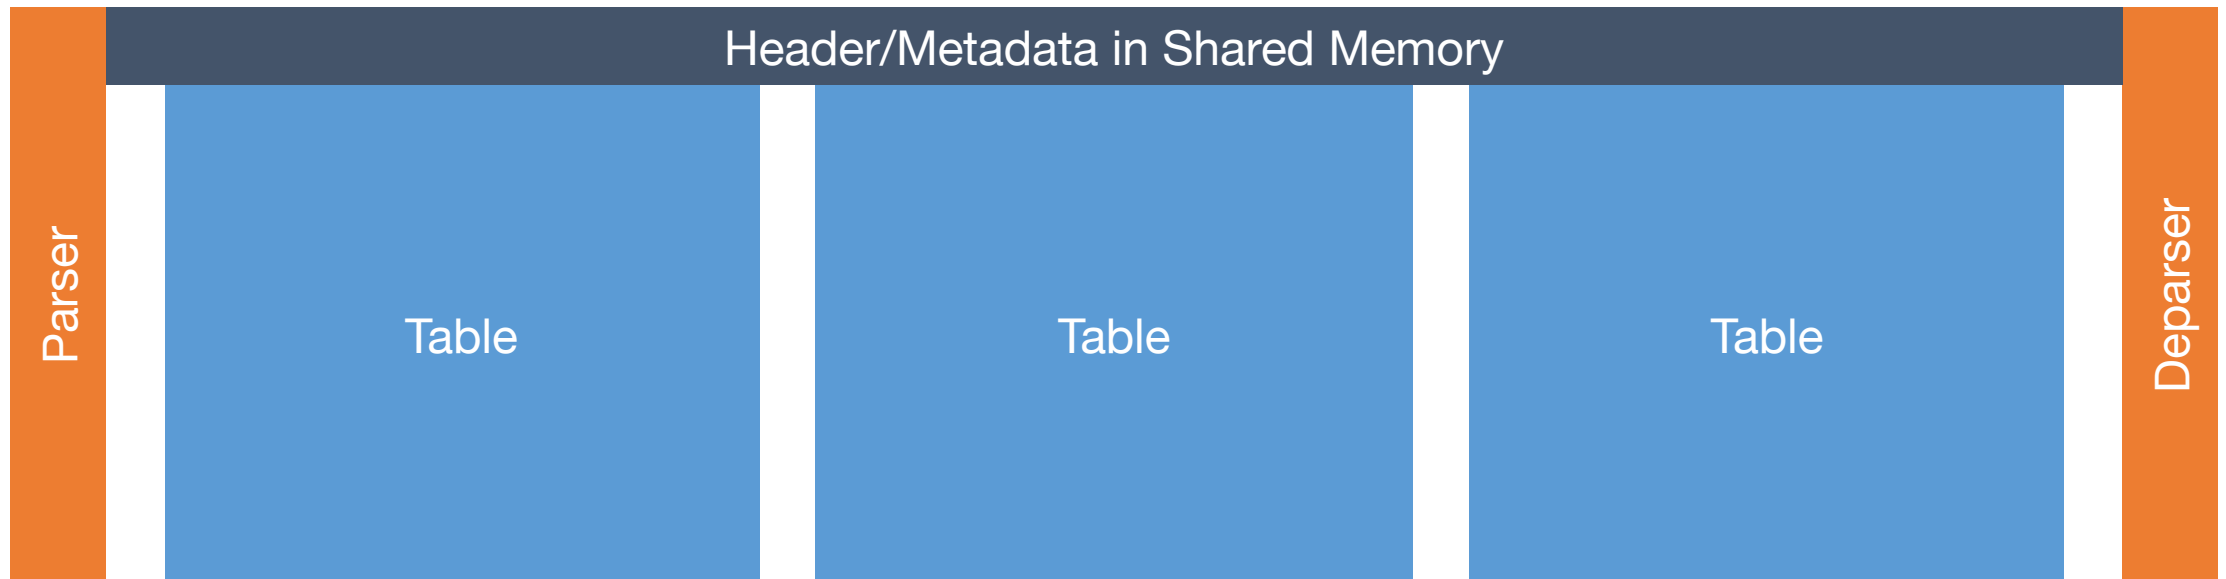
P4: Programmable Protocol-Independent Packet Processing

- Programmable parser: easily add new header
- Programmable table
 - Easily change match fields and actions
 - Easily change rules in the table



P4: Programmable Protocol-Independent Packet Processing

Add new features without changing the hardware



Example: Add IPv6 Support

Parser graph


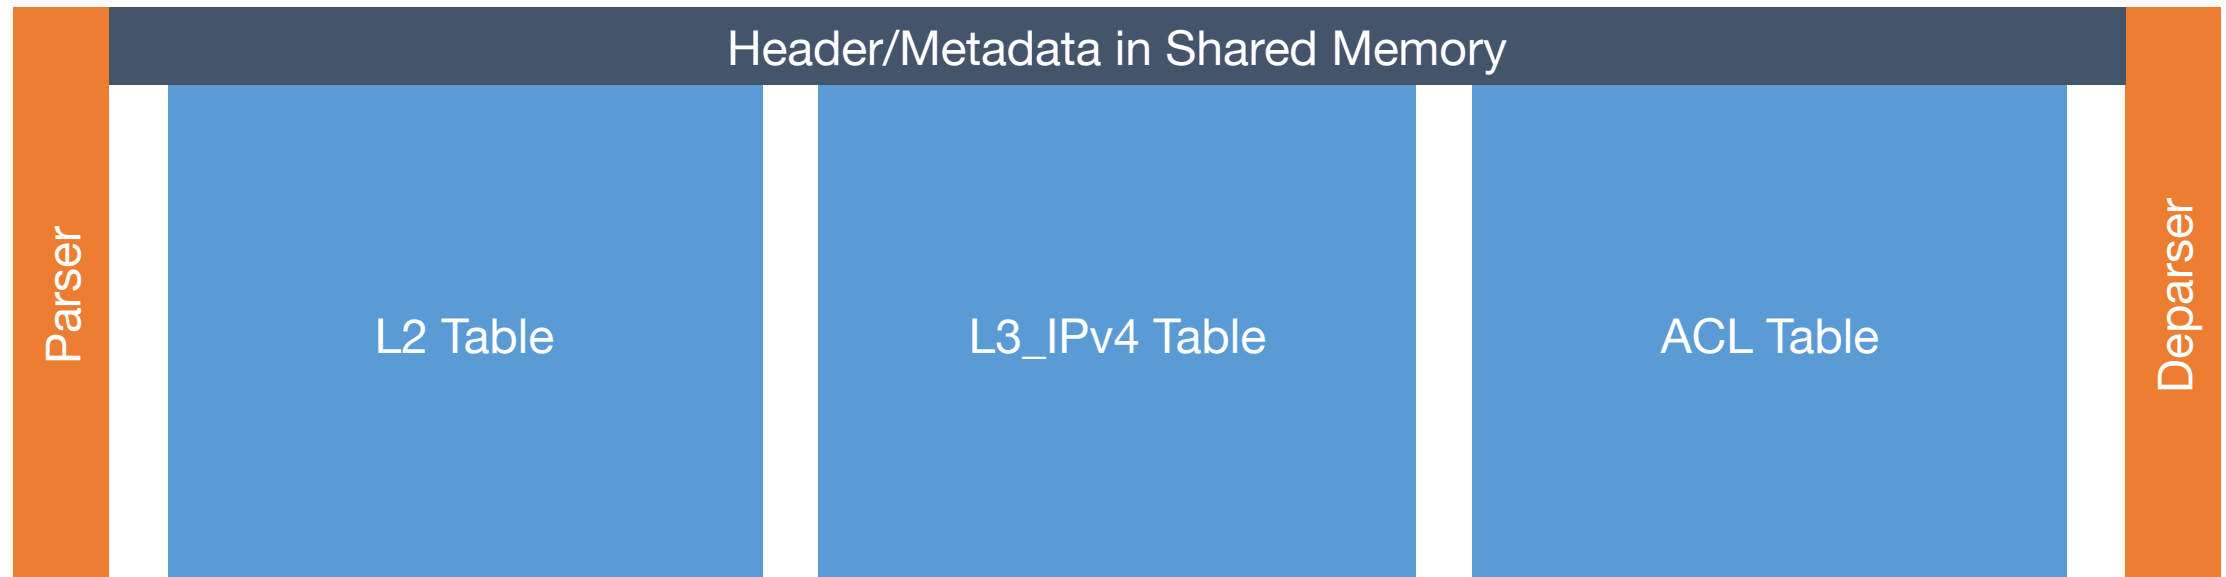



Table graph



Example: Add IPv6 Support

Parser graph


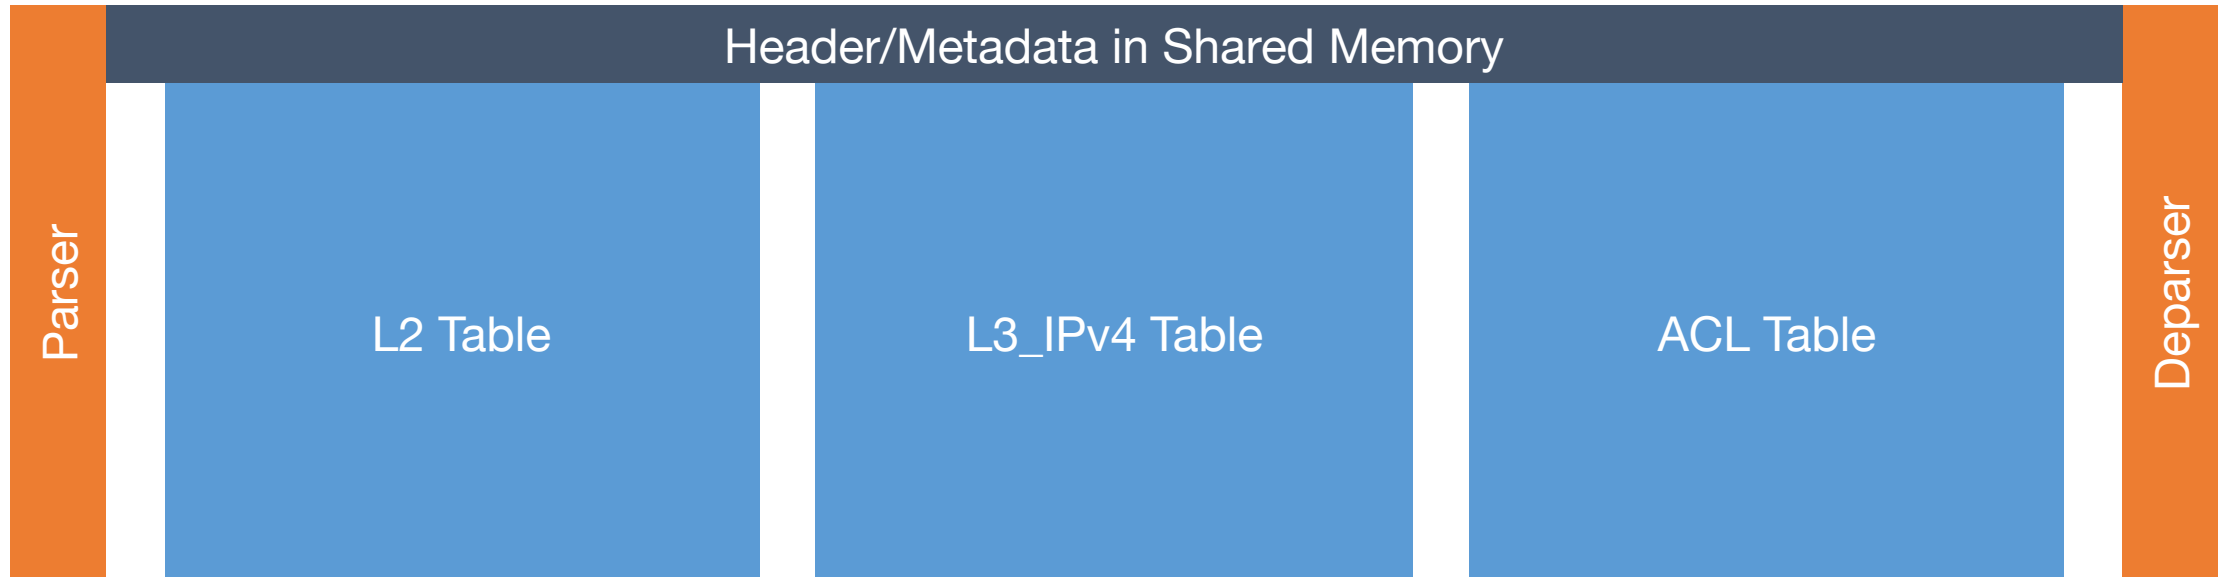



Table graph



Example: Add IPv6 Support

Parser graph


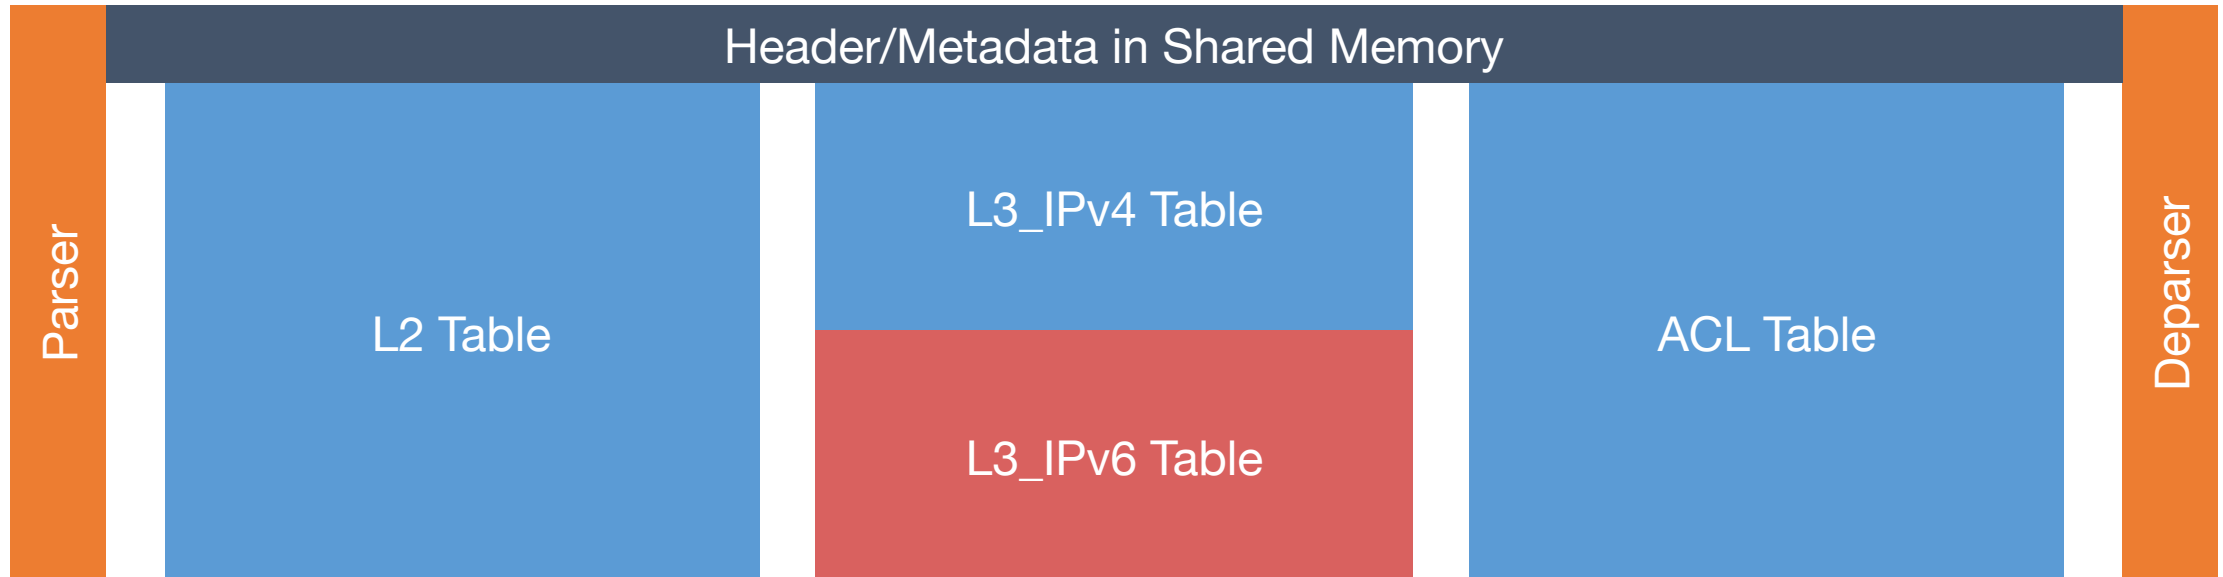





Table graph



P4 Workflow



➤ You need to write

my_switch.p4

Define your switch

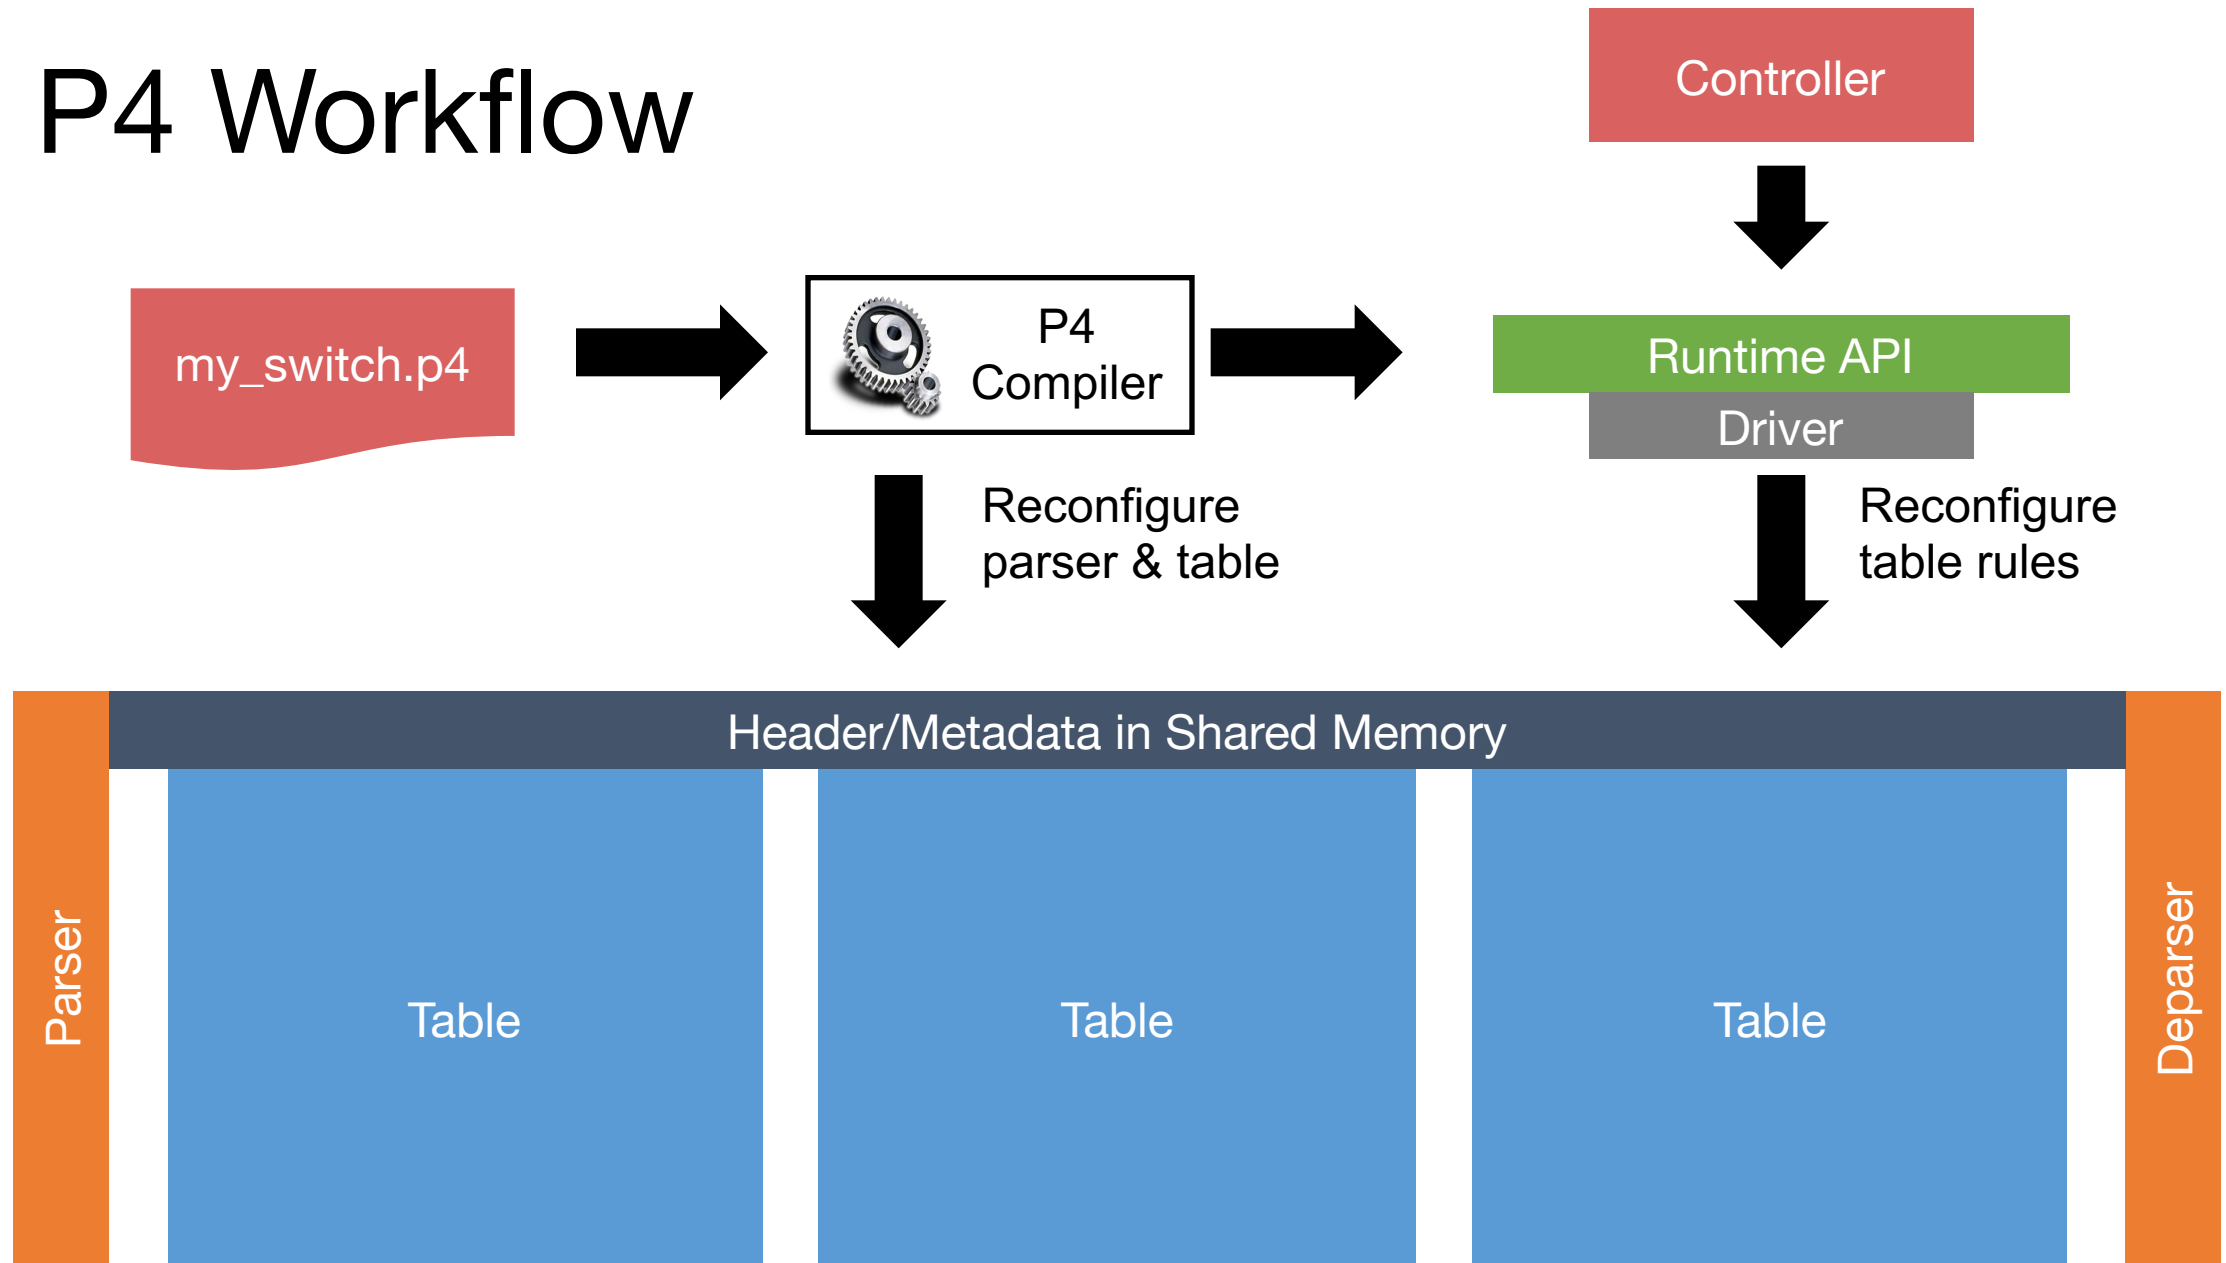
- Specify header format
- Specify parser graph
- Specify type (match & action) of each table
- Specify table graph

Controller

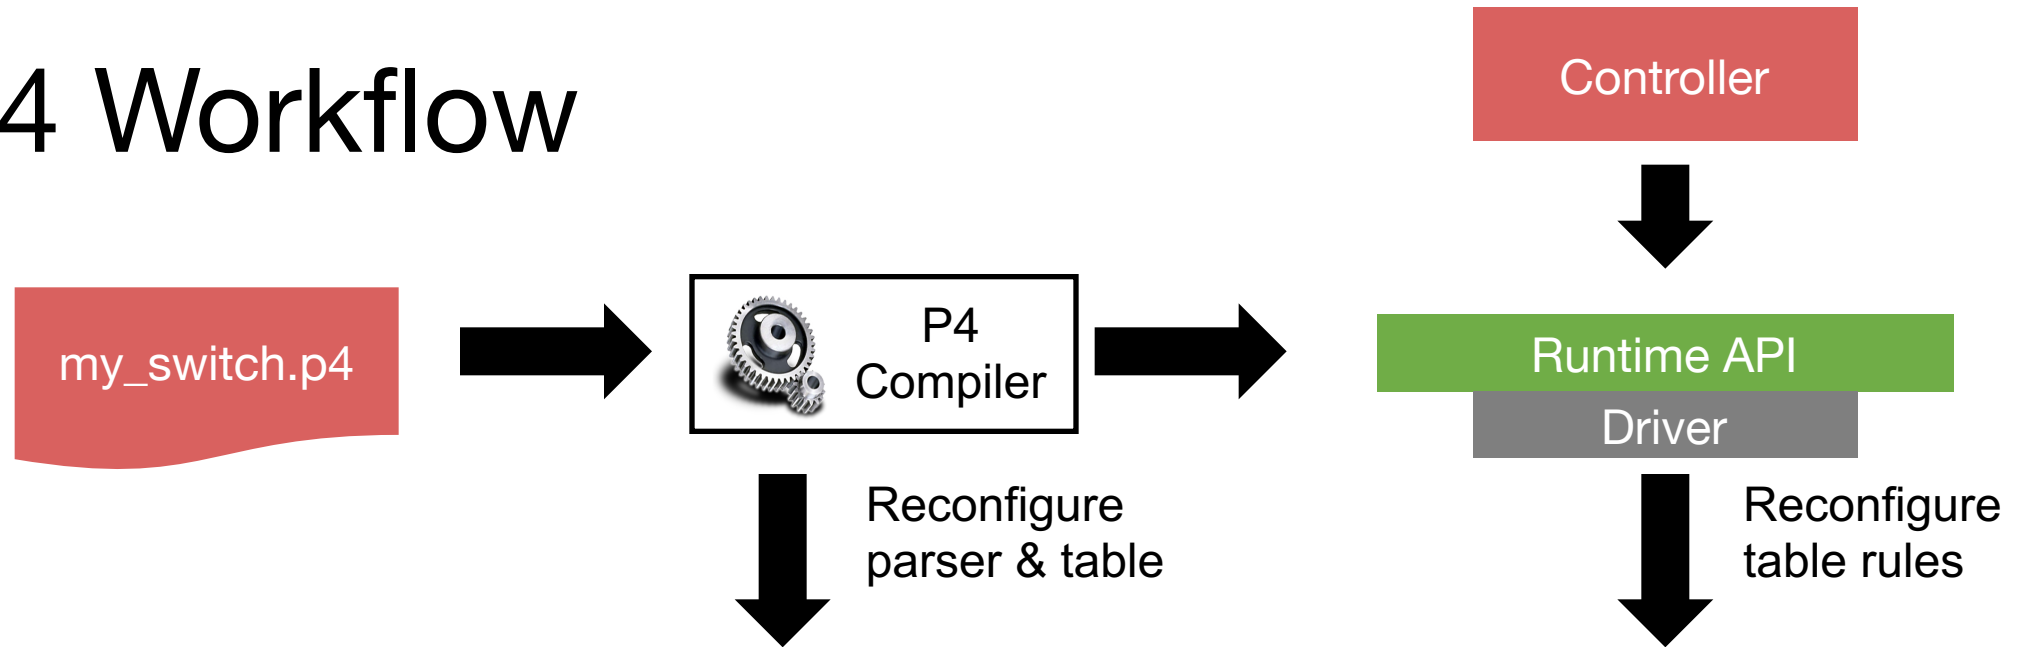
Define how to operate your switch

- Collect traffic stats from switch
- Make control decisions
- Reconfigure rules for each table for each switch

P4 Workflow



P4 Workflow



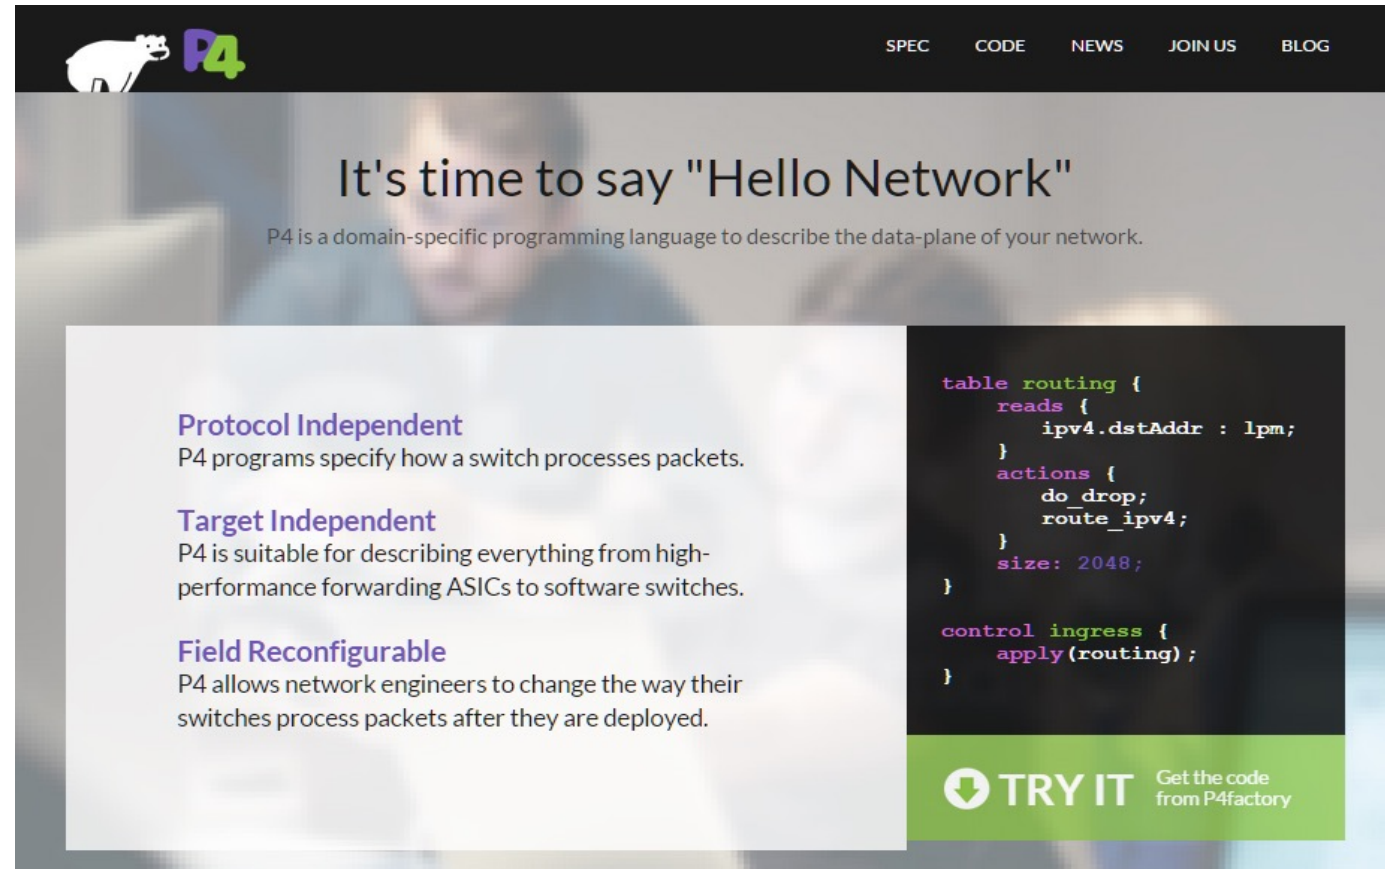
P4 Target

- ASIC
 - Intel Flexpipe, Cisco Doppler, Cavium (Xpliant), Barefoot Tofino, ...
- NPU
 - EZchip, Netronome, ...
- CPU (Software Switch, NIC)
 - Open Vswitch, eBPF, DPDK, VPP...
- FPGA
 - Xilinx, Altera, ...

The P4 Language Consortium

(<http://p4.org>)

- Consortium of academic and industry members
- Open source, evolving, domain-specific language
- Permissive Apache license, code on GitHub today
- Membership is free: contributions are welcome
- Independent, set up as a California nonprofit





P4.org Membership



Original P4 Paper Authors:



Operators



Systems



Targets



Academia



- **Open source**, evolving, domain-specific language
- Permissive Apache license, code on GitHub today

- **Membership is free**: contributions are welcome
- Independent, set up as a California nonprofit

Let's write your first switch in P4...

P4 Workflow

➤ You need to write

my_switch.p4

Define your switch

- Specify header format
- Specify parser graph
- Specify type (match & action) of each table
- Specify table graph

Controller

Define how to operate your switch


- Collect traffic stats from switch
- Make control decisions
- Reconfigure rules for each table for each switch

Write P4 Switch Specification

Header Declaration

my_switch.p4

Packet format



Header/Metadata in Shared Memory

Parser

Table

Table

Table

Deparser

ETH	IPv4	TCP/ UDP	Payload
-----	------	-------------	---------

Write P4 Switch Specification

Header Declaration


my_switch.p4

```
header_type ethernet_t {
  fields {
    dstAddr: 48;
    srcAddr: 48;
    etherType: 16;
  }
}
header ethernet_t ethernet;
```

```
header_type ip4_t {...}
header ipv4_t ipv4;

header_type tcp_t {...}
header tcp_t tcp;

header_type udp_t {...}
header udp_t udp4;
```




Write P4 Switch Specification

```
header_type ethernet_t {  
  fields {  
    dstAddr: 48;  
    srcAddr: 48;  
    etherType: 16;  
  }  
} header ethernet_t ethernet;  
  
header_type ipv4_t (...)  
header ipv4_t ipv4;  
  
header_type tcp_t (...)  
header tcp_t tcp;  
  
header_type udp_t (...)  
header udp_t udp;
```

Parser Graph Declaration

my_switch.p4

Parser graph



Header/Metadata in Shared Memory

Parser


Table

Table

Table

Deparser

Write P4 Switch Specification




```
header_type ethernet_t {  
  fields {  
    dstAddr: 48;  
    srcAddr: 48;  
    etherType: 16;  
  }  
} header ethernet_t ethernet;  
  
header_type ipv4_t {...}  
header ipv4_t ipv4;  
  
header_type tcp_t {...}  
header tcp_t tcp;  
  
header_type udp_t {...}  
header udp_t udp;
```

Parser Graph Declaration


my_switch.p4

```
parser start {  
  extract(ethernet);  
  return select(latest.etherType) {  
    0x0800: parse_ipv4;  
    default: ingress;  
  }  
}
```

```
parser parse_ipv4 {  
  extract(ipv4);  
  return select(latest.protocol) {  
    6:      parse_tcp;  
    17:     parse_udp;  
    default: ingress;  
  }  
}  
parser parse_tcp {...}  
parser parse_udp {...}
```



Write P4 Switch Specification



```
header_type ethernet_t {  
  fields {  
    dstAddr: 48;  
    srcAddr: 48;  
    etherType: 16;  
  }  
}  
header ethernet_t ethernet;  
  
header_type ipv4_t (...)  
header ipv4_t ipv4;  
header_type tcp_t (...)  
header tcp_t tcp;  
header_type udp_t (...)  
header udp_t udp;
```

my_switch.p4

```
parser start {  
  extract(ethernet);  
  return select(latest.etherType) {  
    0x0800: parse_ipv4;  
    default: ingress;  
  }  
}
```

```
parser parse_ipv4 {  
  extract(ipv4);  
  return select(latest.protocol) {  
    6:      parse_tcp;  
    17:     parse_udp;  
    default: ingress;  
  }  
}  
parser parse_tcp {...}  
parser parse_udp {...}
```

Header/Metadata in Shared Memory

Parser

Table

Table

Table

Deparser

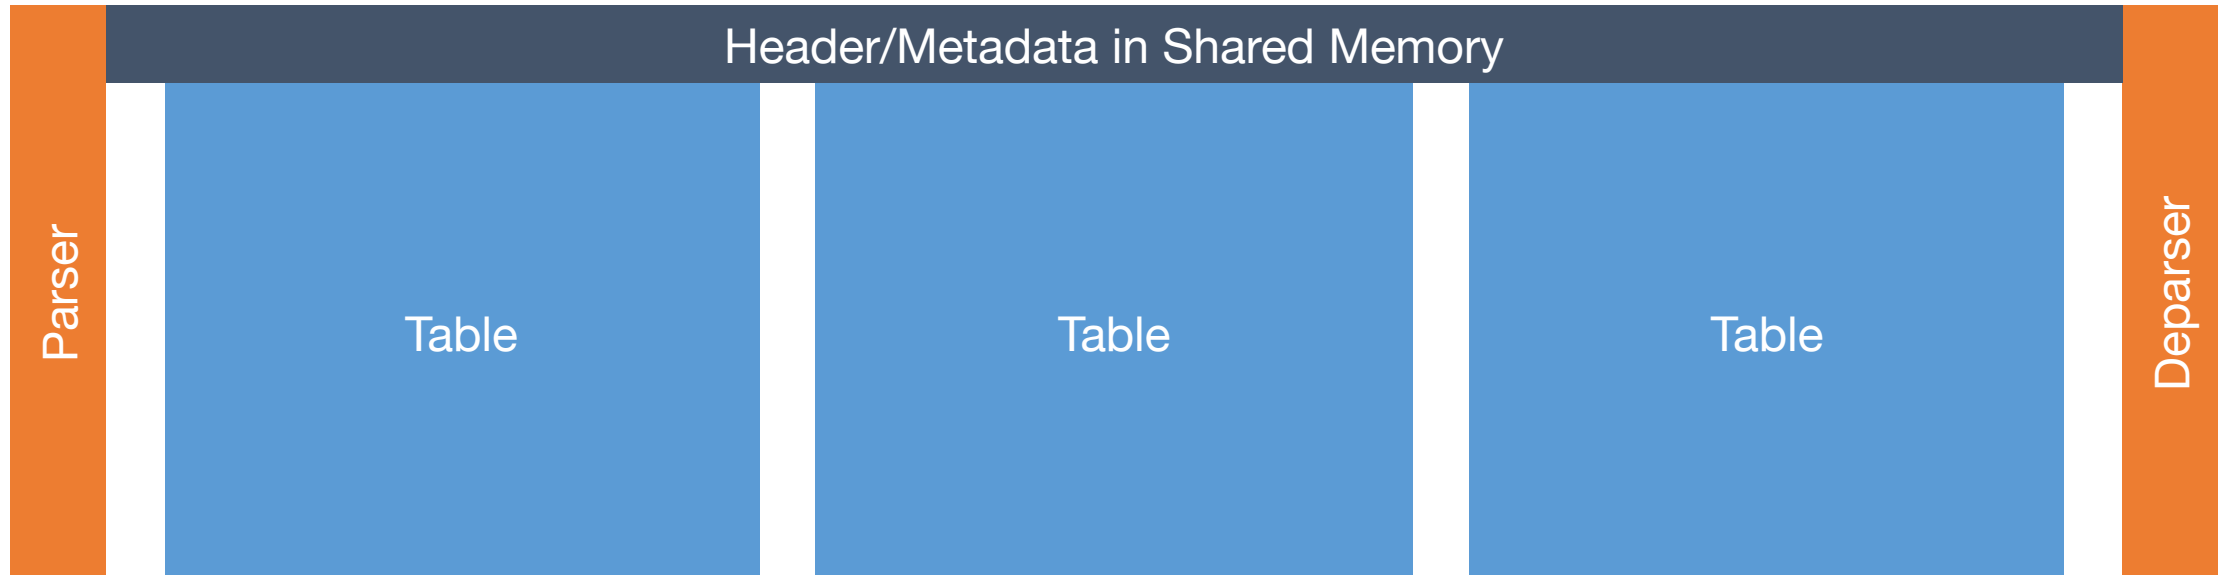
Write P4 Switch Specification

```
header_type ethernet_t {  
  fields {  
    dstAddr: 48;  
    srcAddr: 48;  
    etherType: 16;  
  }  
  header ethernet_t ethernet;  
}  
  
header_type ipv4_t (...)  
header ipv4_t ipv4;  
header_type tcp_t (...)  
header tcp_t tcp;  
header_type udp_t (...)  
header udp_t udp;
```

Table Declaration

my_switch.p4

```
table L2 {  
  reads {  
    eth.dstAddr: exact;  
  }  
  actions {  
    set_port;  
    no_op;  
  }  
  size: 4096 }
```



Actions

- Primitive actions
 - no_op, drop
 - modify_field, modify_field_with_hash_based_offset
 - add, add_to_field
 - register_read, register_write
 - add_header, remove_header, copy_header, push/pop (a header)
 - count, execute_meter
 - resubmit, recirculate, clone
 - etc.

- Compound actions

```
action compound_action(dst_port, dst_ip) {  
    modify_field(standard_metadata.egress_spec, dst_port)  
    modify_field(ipv4.dst_ip, dst_ip);  
    add_to_field(ipv4.ttl, -1)  
}
```

Write P4 Switch Specification

```
header_type ethernet_t {  
  fields {  
    dstAddr: 48;  
    srcAddr: 48;  
    etherType: 16;  
  }  
  header ethernet_t ethernet;  
}  
  
parser start {  
  extract(ethernet);  
  return select(test.etherType) {  
    0x0800: parse_ipv4;  
    default: ingress;  
  }  
}  
  
header_type ipv4_t {...}  
header ipv4_t ipv4;  
  
header_type tcp_t {...}  
header tcp_t tcp;  
  
header_type udp_t {...}  
header udp_t udp;
```

Table Declaration

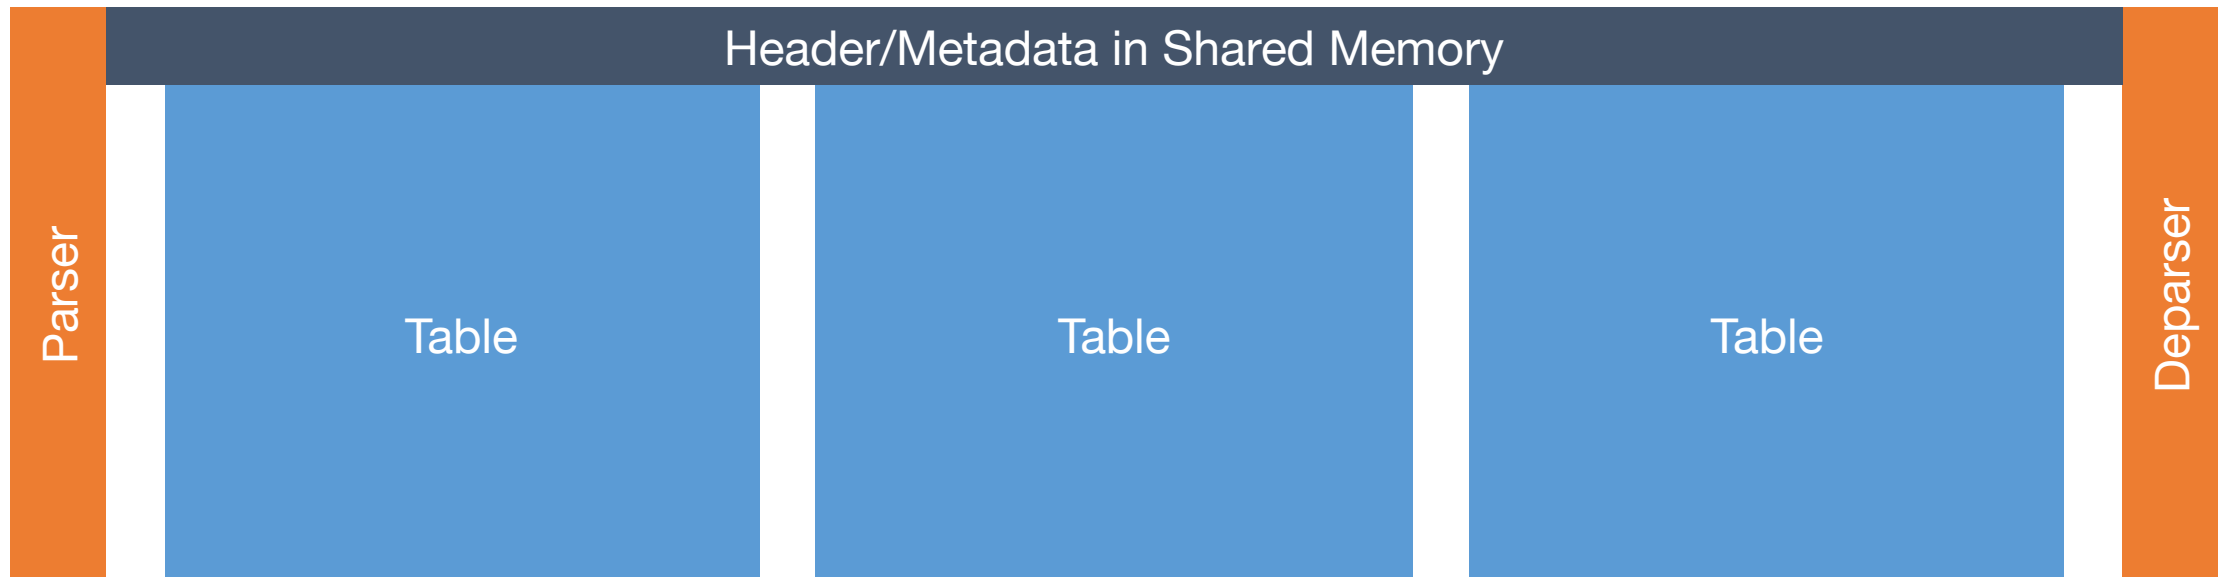
my_switch.p4

```
table L2 {  
  reads {  
    eth.dstAddr: exact;  
  }  
  actions {  
    no_op;  
    set_port;  
  }  
  size: 4096 }
```

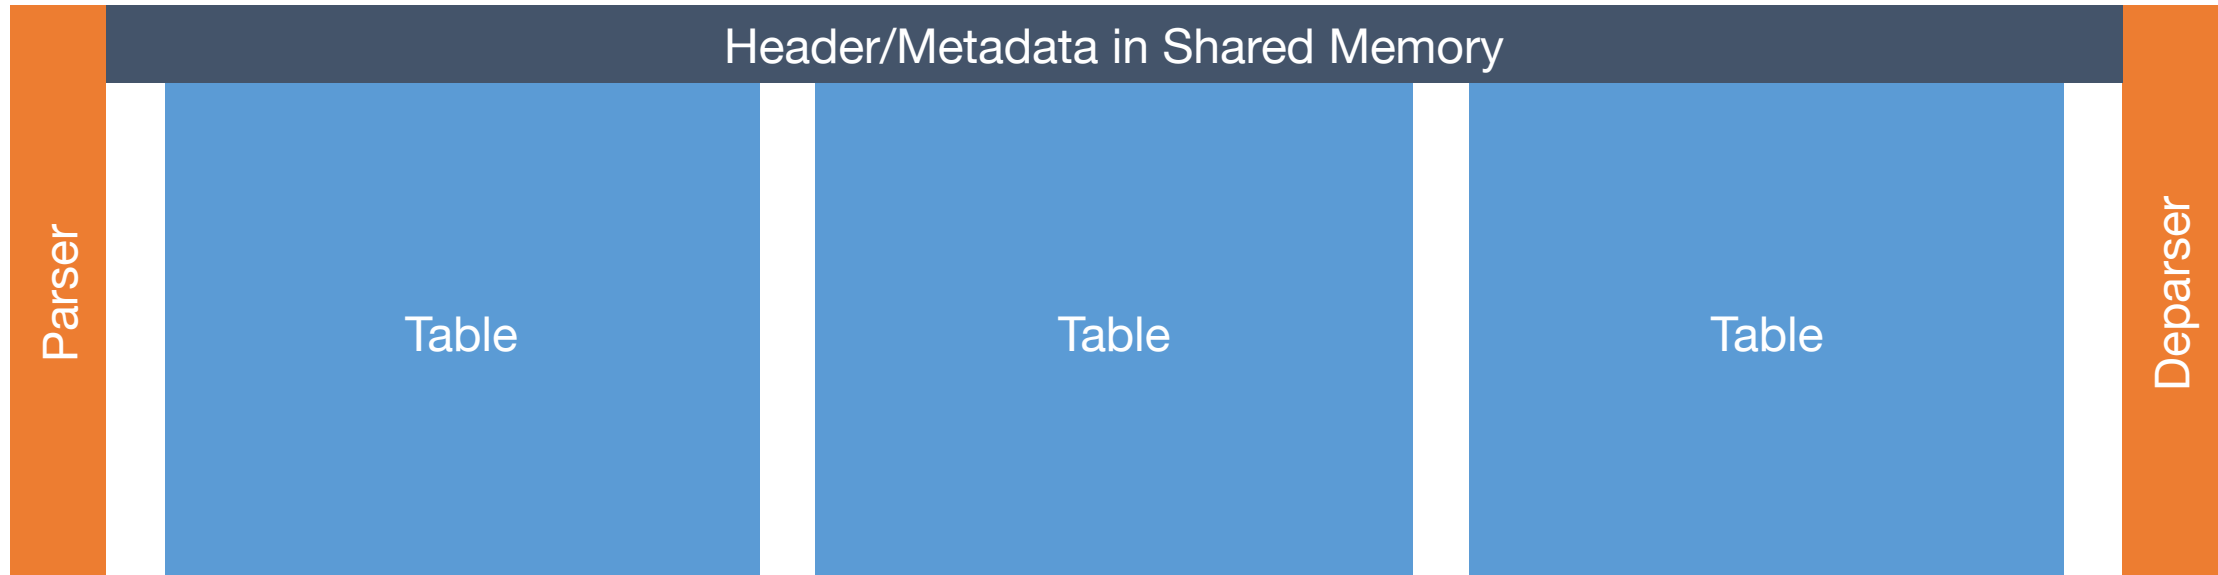

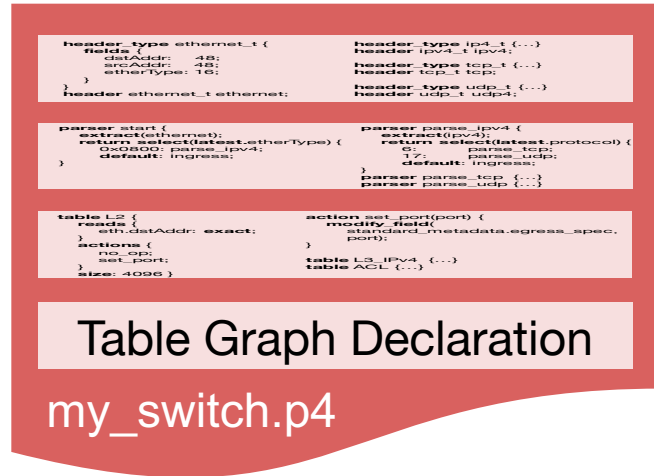
```
action set_port(port) {  
  modify_field(  
    standard_metadata.egress_spec,  
    port);  
}
```

table L3_IPv4 {...}

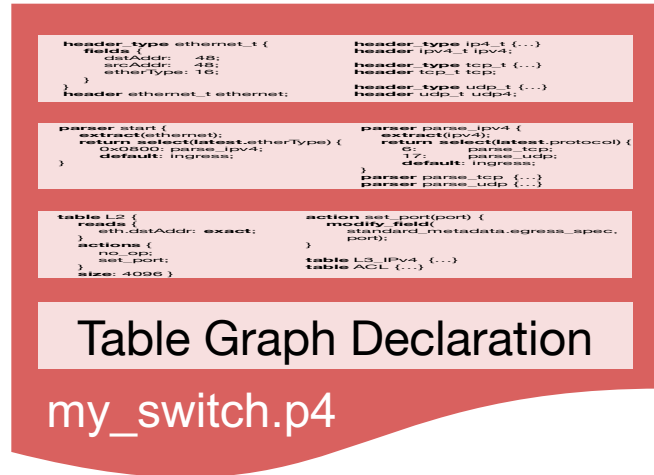

table ACL {...}




Write P4 Switch Specification



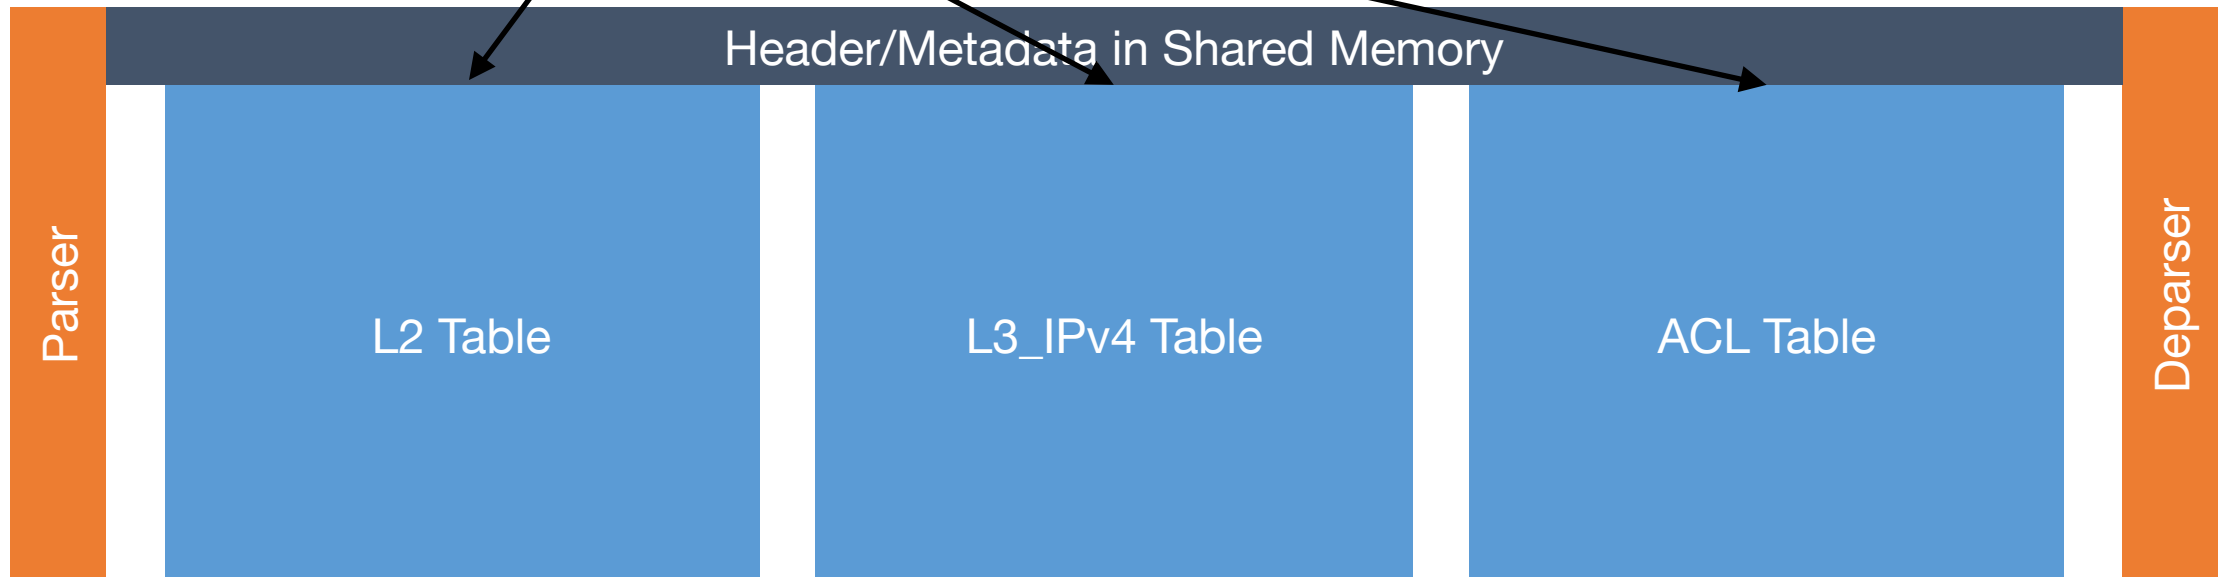
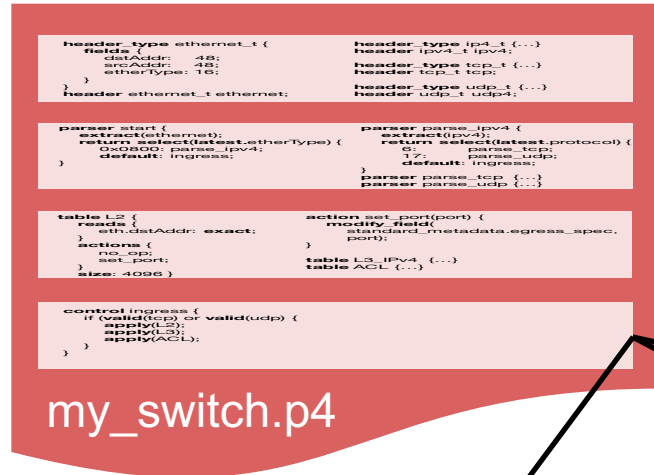

Write P4 Switch Specification



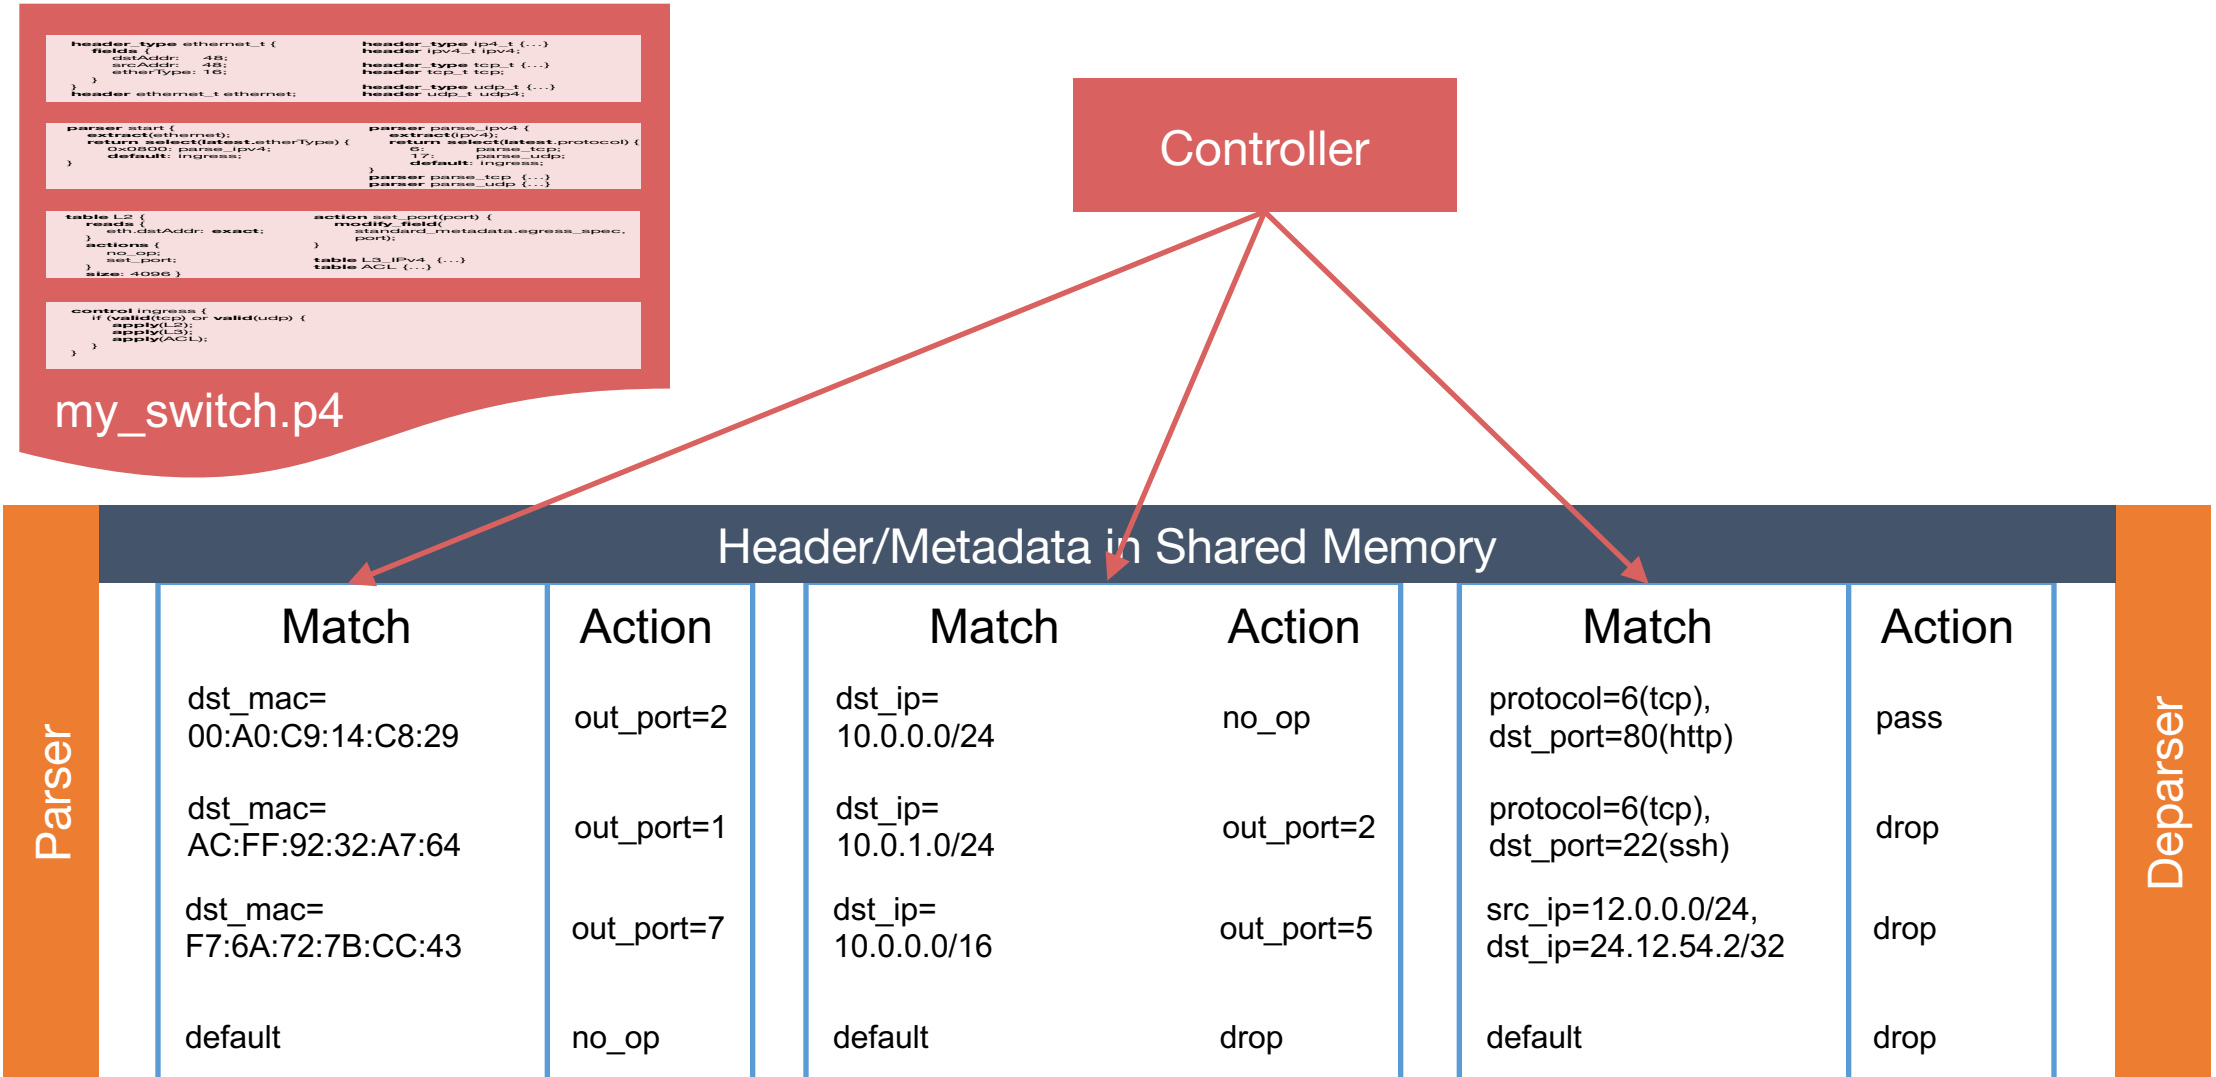
```
control ingress {
  if (valid(tcp) or valid(udp) {
    apply(L2);
    apply(L3);
    apply(ACL);
  }
}
```



Write P4 Switch Specification



Write Controller to Install Rules



Thanks!